Supplementary information
Few-label calibration under publication shift in catalysis

1 Analysis populations and provenance

Historical cleaned input snapshots are inherited from the supplied research archive; the preceding full-study release refit the controlled split analysis and established a traceable source-tuned adaptation engine. It does not claim independent re-extraction of every historical measurement or restoration of the original archive's missing executed-engine version. Rows preserve original-file and row identifiers; historical OCM includes the row-audit table. The cleaned snapshots, their hashes and actual train/test ledgers are the reproducible starting point. Historical raw-file curation exclusions are not represented as newly independently audited extraction decisions.

The methanol paired corpus has 1,245 rows, 82 sources and 80 leakage groups; STY has 1,230 rows and 109 sources; OCM has 1,794 rows, 408 sources and 405 groups. Methanol and STY share 20 publication identifiers. OCM F0/F1 have identical outcome and source populations. All reported targets are finite; genuine zeros remain observations. Required operating-variable meanings follow the released upstream datasets. Missing operating values are imputed using a median fitted on the current training rows; an entirely missing training column uses the documented keep-empty-feature behavior. Numeric scaling also uses only that partition. No target-source query feature values extend composition vocabularies.

Table 1: Historical input quality, recomputed from distributed snapshots.

Panel	Feature	Missing	Minimum	Maximum
methanol	Temperature_K	0	433	673
methanol	Pressure_bar	0	0.4	80
methanol	H2_CO2_ratio	0	1	10
methanol	GHSV_L_g_h	77	0.87	240
methanol	BET_surface_area_m2_g	230	2	974.9
ocm_f0	Temperature_K	0	773	1223
ocm_f0	CH4_partial_pressure_bar	0	0.0094	2.7
ocm_f0	O2_partial_pressure_bar	0	0.003	1.35
ocm_f0	CH4_O2_ratio	0	0.8	38.31
ocm_f0	total_pressure_bar	0	0.106	4.5
ocm_f0	contact_time_s	44	0.0017	182.5
sty	Temperature_K	0	433	673
sty	Pressure_bar	0	1	360
sty	H2_CO2_ratio	0	1	20
sty	GHSV_L_g_h	0	0.9	120
sty	SBET_m2_g	0	2	1102

There are 497 missing catalyst names in the paired methanol data. The name holdout uses only nonempty normalized strings and requires a complete eligible stratum; it does not identify a physical preparation batch. The STY operating snapshot has no catalyst-name field. Richer public Suvarna metadata are retained in the retrieved workbook but are not retroactively added to this operating-input benchmark. The OCM catalyst number identifies an upstream record, not repeated measurements certified on one physical catalyst.

2 Chronology and claim scope

The historical adaptation study is retrospective: archived outcomes motivated its research question. The new-source protocol was locally timestamped and hashed before new query outcomes were evaluated, not independently preregistered. Its primary hypothesis is the five-label composition-enriched blend versus local response contrast. Bibliographic identity corrections remove re-listed old sources even when spelling or number changes; no error-dependent inclusion is used. The additional compilation is a separately released source of observations, rather than a prospective laboratory campaign.

The label-permutation, matched-alpha and fit-regime sensitivity analyses are explanatory follow-ups, not replacements for the frozen primary test. The latter uses every one of the 25 historical outer-training fits; no fit is selected by external performance. Fixed ridge/boosting split audits on old data are retrospective, with configuration held constant within each comparison. Published-workflow audits serve a separate question about validation scope. Suvarna numbers are recomputed from archived predictions; Xu's matched five-fold grouping comparison is newly refitted with XGBoost 3.2.0. Neither is external validation of the new adapters.

3 Split and support construction

Historical outer folds are stored in leakage-group split ledgers, with seeds 42, 101, 202, 303 and 404. Each seed covers every row exactly once as an outer test observation. Training inner folds randomly permute sorted leakage-group identities into three blocks. All imputation/scaling/composition encoders are fitted on the relevant inner training block before OOF prediction.

For a source with n records, sort row identifiers, shuffle with a stable SHA-derived identifier seed, and query $\max(5, \lceil n/4 \rceil)$ rows. Remaining rows are the support pool. A complete curve requires ten pool records. Fifty independently seeded pool orders create nested budgets 1, 2, 3, 5 and 10; the same order is used for every method. No support record is also a query. Calibration labels are disclosed to the prediction function only for that budget.

High-temperature design enumerates observed cutoffs and selects the upper block closest to a quarter of available finite-temperature records, subject to at least five queries and ten strictly lower-temperature pool rows. It never splits tied temperatures across support and query. Name-stratum design enumerates complete reported-name strata with the same minimum sizes and selects the query size closest to a quarter; it never places the same normalized name on both sides. Stress hyperparameters are frozen from random-source training. Matched-source random panels are available for interpretation, but their query records differ from stress panels, so these contrasts are not causal measurements of a difficulty increase.

4 All nine strategies

Let $p_A = f_G(A)$, $p_T = f_G(T)$, support responses y_A, support residual mean $\bar{r}$, and historical response scale $s > 0$.

1. Global: $q = p_T$, a zero-support-label reference at every plotted budget.
2. Support mean: $q = \bar{y}_A$.
3. Offset: $q = p_T + \bar{r}$.
4. Affine: $z_A = (p_A - \bar{p}_A)/s$, $z_T = (p_T - \bar{p}_A)/s$,

$$q = p_T + \bar{r} + z_T \frac{z_A^\top (y_A - p_A)}{z_A^\top z_A + \lambda}.$$

This shrinks the additional slope toward an offset map, with an unpenalized residual mean.

5. Empirical-Bayes offset: $q = p_T + \mu + \eta_k(\bar{r} - \mu)$,

$$\eta_k = \frac{k\tau^2}{\sigma^2 + k\tau^2}.$$

Historical OOF residuals provide the source-equal residual mean μ, pooled within-source variance σ^2, and between-source variance corrected for within-source mean noise, truncated below at zero. The variance floor is 10^{-12}.

6. Local ridge: fit y_A on training-fitted z_A, intercept unpenalized.

7. Temperature ridge: the same support fit using only the scaled temperature coordinate.

8. Residual ridge: fit $y_A - p_A$ on z_A and add p_T.

9. Shrinkage blend: combine local and residual predictions with support leave-one-out and historical pseudo-source shrinkage, as specified in the main text.

Ridge/affine strengths are separately selected per method and label budget from $\{.1, 1, 10, 100\}$. Ties follow the declared grid order. The base grid has four ridge candidates; boosting candidates (100,2), (100,3), (300,2), (300,3), (500,8), where pairs are trees/depth. Boosting uses learning rate .05, subsample .8, minimum split size 3 and random state 42. Its candidate score is source-equal absolute OOF error, not pooled record MAE.

Adapters are scored on ten training-source pseudo-support draws per budget. Local, residual, temperature and affine strengths minimize source-equal pseudo-query MAE. For blending, prior $w_0 \in \{0, .25, .5, .75, 1\}$ minimizes pseudo-query error; disagreement scale is the source-equal mean component leave-one-out squared difference (one-label query disagreement at $k = 1$). Shrinkage $c \in \{\infty, 10, 2, .5, 0\}$ is selected with the prior/scale for each pseudo-source excluding its leakage group. Actual outer or additional-source query labels are not used. Reuse of training pseudo-queries across hyperparameter stages can affect training selection, but final outer and additional-source assessments remain separate.

5 Exact ridge leave-one-out and projection identity

Center support matrix X and responses y without penalizing the intercept. For k support records, let $K = X_c X_c^\top$, $B = (K + \alpha I)^{-1}$, and

$$\widehat{y}_T = \bar{y} + (X_T - \bar{X}) X_c^\top B(y - \bar{y}).$$

The diagonal leverage is $h_i = 1/k + (KB)_{ii}$. For $k > 1$ the exact leave-one-out response prediction is

$$\widehat{y}_{-i} = y_i - \frac{y_i - \widehat{y}_i}{1 - h_i}.$$

This uses fixed historical transformations; it does not refit the historical preprocessing on support labels. Independent mathematical verification compares predictions with scikit-learn Ridge and leave-one-out refits.

For fixed features, the map from y_A to query predictions is linear. Therefore

$$R_\alpha[y_A - p_A] = R_\alpha[y_A] - R_\alpha[p_A], \quad q_R - q_L = p_T - R_\alpha[p_A].$$

The main method chooses its two alphas separately. In the exploratory matched-alpha ablation, both use the selected local alpha without new tuning. The identity is checked numerically for all historical OCM and additional-source random draws at $k = 5$. An added global component can correct a local response or add error; linearity alone does not determine which.

6 Additional OCM source reconstruction

The public updated workbook has 4,759 rows and 533 represented publication numbers. The old bibliography has 420 unique numbers up to 421, including an ambiguous duplicated number. Existing numbers are conservatively old, even when their citation text has a typographic change. Additional numbers are screened against normalized old citations and bibliometric fields. Three aliases are removed: 424=396, 427=395 and 429=105, leaving 118 additional bibliographic entries before record-quality eligibility. Four old-number citation differences are logged; two concern missing old citation strings and the others typography. They do not qualify as new sources.

Admitted new records require finite positive temperature, partial pressures, pressure ratio and total pressure; contact time may be missing but cannot be nonpositive when observed. C_2 yield must be reported within $[0, 100]$. Composition encoding validates species/amount pairing and amount range; a fifth/sixth cation or third support cannot be omitted to fit the frozen schema. Full upstream record duplicates are removed. The exclusion reasons overlap: 1,845 records are old/unverifiable publications, 393 lack valid CH_4 partial pressure, 393 lack valid O_2 partial pressure, 100 exceed slots, 44 lack a valid target, and 21 are full duplicates. Their sum is not the excluded-row count. Exactly 2,378 distinct rows are excluded, leaving 2,381 from 91 sources.

The workbook describes species amounts as molar percentages, pressures in bar, temperature in K, contact time in s and C_2 yield as methane conversion times combined C_2 selectivity. We preserve these reported variables rather than infer pressure assumptions or physical absence from blank slots. The complete row-audit, all-source registry and bibliography screen are distributed. The 52-source main cohort and 40-source high-temperature cohort have label-blind design ledgers and shared F0/F1 support orders.

The full-fit F1 backbone is gradient boosting (500,8). At $k = 5$, local and residual alphas are both 1, affine strength 1, temperature alpha 1, blend prior .5, shrinkage 10 and disagreement scale 6.6618. F0 backbone is (100,3); local alpha 10, residual alpha 100, affine .1, temperature 1, blend prior 0, shrinkage 10 and scale 1.4613. These settings come exclusively from old OCM sources. Serialized backbones, preprocessing parameters, training OOF records and all selection scores are saved.

Table 2: Additional OCM main-cohort sources. No. is the public update bibliography number; n is the number of admitted records. Errors are pp at k=5.

No.	Citation	n	Local	Blend	Δ
422	Appl. Catal. A Gen. 219 (2001) 61-68	37	2.927	3.041	0.114
423	Appl. Catal. A Gen. 354 (2009) 143–152	60	0.570	0.753	0.183
425	Ind. Eng. Chem. Res. 48 (2009) 641–646	65	1.853	1.670	-0.183
431	J. Nat. Gas Chem. 19 (2010) 35–42	20	3.420	2.489	-0.931
434	J. Nat. Gas Chem. 19 (2010) 534–538	20	3.302	2.041	-1.260
435	J. Nat. Gas Chem. 19 (2010) 25–30	30	1.773	1.821	0.048
439	Appl. Catal. A Gen. 375 (2010) 172–178	52	2.133	2.102	-0.031
444	J. Nat. Gas Chem. 20 (2011) 204–213	45	7.283	7.467	0.184
445	J. Nat. Gas Chem. 20 (2011) 428–434	63	4.575	3.853	-0.722
446	J. Mol. Catal. A Chem. 346 (2011) 46–54	52	2.196	2.330	0.134
447	Appl. Catal. A Gen. 406 (2011) 1–12	54	4.498	4.643	0.145
448	Appl. Catal. A Gen. 391 (2011) 205–214	28	3.285	3.223	-0.063
449	J. Catal. 281 (2011) 241–253	19	4.873	4.756	-0.118
451	J. Nat. Gas Chem. 21 (2012) 49–55	27	9.719	8.658	-1.061
455	Asian Journal of Chemistry 25 (2013) 9179-9184	35	3.009	3.189	0.180
457	Appl. Catal. A Gen. 454 (2013) 100–114	50	1.466	1.250	-0.217
466	Appl. Catal. A Gen. 485 (2014) 10–19	36	4.163	4.506	0.343
467	ACS Catal. 4 (2014) 1972–1990	62	4.062	4.185	0.122
468	Appl. Catal. A Gen. 484 (2014) 97–107	25	4.732	4.446	-0.286
469	Inorg. Mater. 50 (2014) 912–916	17	3.849	4.261	0.413

No.	Citation	n	Local	Blend	Δ
470	Appl. Catal. A Gen. 472 (2014) 39–46	23	5.613	5.432	-0.182
471	Dokl. Phys. Chem. 455 (2014) 60–63	18	3.294	3.592	0.298
472	Catal. Today 228 (2014) 5–14	67	1.760	2.292	0.532
476	React. Kinet. Mech. Catal. 112 (2014) 227–240	21	4.086	4.135	0.049
480	Appl. Catal. B Environ. 150–151 (2014) 496–505	21	9.759	9.728	-0.031
481	Nanoscale 7 (2015) 2260–2264	87	5.825	5.945	0.120
482	J. Catal. 329, (2015), 560–573	26	0.432	0.479	0.048
485	ACS Catal. 5 (2015) 1663–1674	20	7.073	7.024	-0.049
486	Pet. Chem. 55 (2015) 163–168	39	3.438	3.481	0.043
492	Appl. Catal. A Gen. 528 (2016) 175–190	140	3.912	3.910	-0.002
496	RSC Adv. 6 (2016) 34872–34876	32	1.246	1.296	0.050
503	Appl. Catal. A Gen. 537 (2017) 33–39	57	1.883	1.971	0.088
504	Sci. Adv. 3, (2017)	47	3.181	3.022	-0.159
506	Appl. Catal. A Gen. 544 (2017) 77–83	48	2.532	2.556	0.024
508	Mol. Catal. 435 (2017) 13–23	19	2.221	2.316	0.095
510	ACS Comb. Sci. 19 (2017) 15–24	32	1.262	1.396	0.135
514	ChemNanoMat 4 (2018) 487–495	44	5.470	5.647	0.177
515	Appl. Catal. A Gen. 552 (2018) 117–128	31	6.306	5.840	-0.466
516	Catal. Today 299 (2018) 60–66	23	4.515	4.592	0.078
520	Chinese J. Catal. 39 (2018) 1395–1402	30	5.390	5.383	-0.007
521	Front. Chem. 6 (2018) 1–10	36	2.325	2.354	0.029
522	Eur. J. Inorg. Chem. (2018) 1787–1799	40	1.853	1.924	0.071
523	Chem. Eng. Sci. 175 (2018) 199–207	22	2.275	2.421	0.146
526	Mol. Catal. 457 (2018) 74–81	27	1.930	2.000	0.070
527	Inorg. Chem. Commun. 100 (2019) 97–100	28	5.725	5.490	-0.235
528	ACS Catal. 9 (2019) 4030–4045	18	3.406	3.711	0.305
529	Mol. Catal. 470 (2019) 40–47	146	1.595	1.796	0.201
531	J. Catal. 375 (2019) 478–492	26	4.362	4.578	0.217
535	Eur. J. Inorg. Chem. 2 (2019) 183–194	30	3.588	3.385	-0.203
536	Polyhedron 170 (2019) 602–611	63	1.200	1.175	-0.024
537	Chinese Chem. Lett. 30 (2019) 1141–1146	25	2.919	2.318	-0.601
539	Ind. Eng. Chem. Res. 58 (2019) 22847–22856	24	2.746	2.803	0.057

7 Metric estimands and uncertainty

Let e_{srj} be query MAE for source s, support repeat r and historical outer seed j. Historical source MAE is $\bar{e}_s = \frac{1}{J}\sum_j \frac{1}{R}\sum_r e_{srj}$, and the primary benchmark MAE is $\frac{1}{S}\sum_s \bar{e}_s$. External full-fit MAE omits the seed average. Row-weighted sensitivity weights each source by its fixed query count. Bounded sensitivity clips percentage predictions to $[0, 100]$ and STY predictions below at zero; raw observations are unchanged.

Cluster bootstrap samples leakage groups with replacement and divides sampled source-error sums by sampled source counts. It does not replace the source-equal estimand with group-equal averaging. Twenty thousand draws estimate percentile intervals, conditional on the fitted models and observed sources. No adjustment across secondary comparisons is applied. A zero-crossing interval is neither equivalence nor proof of no effect. Identical within-source control outputs can yield numerical differences near floating-point precision.

Stacked repeated R^2 is $1 - \sum_{sjri}(y_i - p_i)^2 / \sum_{sjri}(y_i - \bar{y})^2$ over the specified repeated prediction population. R^2 of a mean prediction instead averages each record's predictions before squaring the error; these are different quantities. The main controlled audit uses the former. An older archive field named prediction-averaged R^2 actually averaged squared errors before forming R^2; that label is not retained as an averaged-prediction metric.

8 Complete budget tables

Values below are source-equal raw MAEs. Units are pp for percent targets and mg MeOH h^{-1} g_{cat}^{-1} for STY. Stress cohorts and selected parameters are fixed across budgets. Intervals, bounded/row-weighted sensitivities and paired source values accompany the CSV exports; printed tables avoid pretending that support repeats increase the number of independent sources.

Historical / methanol / CO2_conversion_percent / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.271	5.271	5.271	5.271	5.271
Support mean	5.309	4.616	4.420	4.304	4.146
Offset	4.269	3.727	3.556	3.455	3.329
Affine	4.269	3.635	3.336	3.083	2.810
EB offset	4.044	3.699	3.551	3.457	3.331
Local ridge	5.309	3.633	2.914	2.429	2.039
Temperature ridge	5.309	3.526	2.905	2.519	2.267
Residual ridge	4.269	3.402	3.009	2.720	2.368
Shrinkage blend	4.315	3.394	2.870	2.441	2.040

Historical / methanol / CO2_conversion_percent / stratum_holdout

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.352	5.352	5.352	5.352	5.352
Support mean	5.832	5.112	4.878	4.673	4.581
Offset	5.081	4.494	4.301	4.155	4.080
Affine	5.081	4.608	4.333	4.167	4.028
EB offset	4.453	4.185	4.086	4.023	4.011
Local ridge	5.832	4.294	3.732	3.366	3.160
Temperature ridge	5.832	4.355	3.772	3.439	3.280
Residual ridge	5.081	4.339	4.005	3.787	3.568
Shrinkage blend	5.080	4.212	3.740	3.435	3.225

Historical / methanol / CO2_conversion_percent / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	7.536	7.536	7.536	7.536	7.536
Support mean	7.960	7.555	7.521	7.379	7.344
Offset	5.688	5.331	5.244	5.109	5.049
Affine	5.688	5.191	4.927	4.664	4.462
EB offset	5.761	5.428	5.322	5.168	5.078
Local ridge	7.960	5.455	4.098	3.210	2.834
Temperature ridge	7.960	5.262	4.099	3.477	3.067
Residual ridge	5.688	4.849	4.416	3.924	3.519
Shrinkage blend	5.818	4.933	4.129	3.275	2.849

Historical / methanol / Methanol_selectivity_percent / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	22.740	22.740	22.740	22.740	22.740
Support mean	21.217	18.676	17.803	17.136	16.680
Offset	19.098	17.130	16.425	15.747	15.395
Affine	19.098	15.873	14.155	12.539	11.502
EB offset	18.865	17.455	16.738	15.988	15.535
Local ridge	21.217	15.905	13.713	11.547	10.249
Temperature ridge	21.217	15.377	13.309	11.535	10.596
Residual ridge	19.098	15.973	14.168	12.077	10.295
Shrinkage blend	19.625	16.118	13.932	11.584	10.137

Historical / methanol / Methanol_selectivity_percent / stratum_holdout

Strategy	k=1	k=2	k=3	k=5	k=10
Global	25.729	25.729	25.729	25.729	25.729
Support mean	20.946	19.026	18.350	17.767	17.491
Offset	18.503	16.895	16.359	15.838	15.514
Affine	18.503	16.482	15.775	14.971	14.339
EB offset	19.161	17.307	16.701	16.085	15.641
Local ridge	20.946	17.205	15.962	14.510	13.911
Temperature ridge	20.946	17.512	16.092	14.900	14.455
Residual ridge	18.503	16.469	15.598	14.717	13.935
Shrinkage blend	18.799	16.496	15.527	14.523	13.809

Historical / methanol / Methanol_selectivity_percent / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	20.817	20.817	20.817	20.817	20.817
Support mean	25.201	23.391	23.329	22.853	22.609
Offset	18.291	16.696	16.482	15.992	15.608
Affine	18.291	16.116	15.544	14.170	12.659
EB offset	16.381	15.559	15.602	15.409	15.278
Local ridge	25.201	18.747	15.272	12.240	10.414
Temperature ridge	25.201	18.725	14.868	12.549	10.896
Residual ridge	18.291	15.776	14.239	12.570	10.355
Shrinkage blend	19.191	16.082	14.399	12.159	9.999

Historical / ocm_f0 / C2_yield_percent / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.302	5.302	5.302	5.302	5.302
Support mean	4.609	4.108	3.875	3.693	3.594
Offset	4.614	4.106	3.862	3.672	3.563
Affine	4.614	4.106	3.871	3.694	3.585
EB offset	4.527	4.164	3.947	3.759	3.618
Local ridge	4.609	4.099	3.867	3.655	3.559
Temperature ridge	4.609	4.113	3.870	3.677	3.536
Residual ridge	4.614	4.109	3.884	3.701	3.582
Shrinkage blend	4.632	4.114	3.878	3.660	3.568

Historical / ocm_f0 / C2_yield_percent / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	3.792	3.792	3.792	3.792	3.792
Support mean	3.949	3.718	3.503	3.372	3.284
Offset	4.224	3.995	3.762	3.650	3.532
Affine	4.224	3.966	3.736	3.641	3.545
EB offset	3.683	3.693	3.582	3.539	3.479
Local ridge	3.949	3.722	3.519	3.381	3.341
Temperature ridge	3.949	3.778	3.681	3.581	3.978
Residual ridge	4.224	4.011	3.799	3.691	3.711
Shrinkage blend	4.145	3.843	3.649	3.481	3.486

Historical / ocm_f1 / C2_yield_percent / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	4.527	4.527	4.527	4.527	4.527
Support mean	4.609	4.108	3.875	3.693	3.594
Offset	4.341	3.823	3.568	3.423	3.299
Affine	4.341	3.793	3.526	3.360	3.296
EB offset	3.927	3.678	3.525	3.415	3.313
Local ridge	4.609	4.080	3.710	3.374	2.839
Temperature ridge	4.609	4.113	3.870	3.677	3.536
Residual ridge	4.341	3.788	3.430	3.109	2.696
Shrinkage blend	4.331	3.773	3.408	3.081	2.633

Historical / ocm_f1 / C2_yield_percent / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	3.473	3.473	3.473	3.473	3.473
Support mean	3.949	3.718	3.503	3.372	3.284
Offset	3.449	3.086	2.957	2.899	2.807
Affine	3.449	3.080	2.941	2.882	2.708
EB offset	2.845	2.716	2.683	2.714	2.702
Local ridge	3.949	3.752	3.525	3.357	3.021
Temperature ridge	3.949	3.778	3.681	3.581	3.978
Residual ridge	3.449	3.111	2.960	2.866	2.567
Shrinkage blend	3.501	3.158	2.943	2.821	2.570

Historical / sty / STY_mgMeOH_per_h_per_gcat / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	139.258	139.258	139.258	139.258	139.258
Support mean	113.283	100.144	95.452	91.993	88.939
Offset	111.411	97.174	92.012	88.265	85.238
Affine	111.411	91.353	79.821	71.225	65.355
EB offset	104.811	94.186	90.131	86.510	84.124
Local ridge	113.283	87.602	75.645	65.210	54.901
Temperature ridge	113.283	100.188	92.917	87.098	72.872
Residual ridge	111.411	89.773	78.984	69.889	59.293
Shrinkage blend	110.732	84.321	71.439	62.608	52.771

Historical / sty / STY_mgMeOH_per_h_per_gcat / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	160.224	160.224	160.224	160.224	160.224
Support mean	154.074	144.778	140.783	139.115	136.853
Offset	139.665	130.528	126.993	124.055	121.817
Affine	139.665	132.316	122.109	113.600	106.997
EB offset	130.711	124.904	122.730	120.884	120.019
Local ridge	154.074	142.114	135.284	133.260	114.010
Temperature ridge	154.074	146.875	142.091	157.467	148.619
Residual ridge	139.665	124.610	115.905	111.808	93.301
Shrinkage blend	148.286	131.385	115.005	115.145	96.205

Additional OCM / ocm_f0 / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.973	5.973	5.973	5.973	5.973
Support mean	5.465	4.907	4.669	4.490	4.343
Offset	5.480	4.926	4.687	4.502	4.353
Affine	5.480	4.902	4.548	4.284	4.057
EB offset	5.201	4.891	4.710	4.545	4.381
Local ridge	5.465	4.743	4.390	4.045	3.718
Temperature ridge	5.465	4.742	4.106	3.779	3.580
Residual ridge	5.480	4.771	4.206	4.421	3.665
Shrinkage blend	5.444	4.736	4.261	4.054	3.632

Additional OCM / ocm_f0 / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.699	5.699	5.699	5.699	5.699
Support mean	6.204	5.886	5.712	5.556	5.436
Offset	6.217	5.843	5.674	5.512	5.384
Affine	6.217	5.813	5.484	5.209	4.946
EB offset	5.183	5.219	5.219	5.205	5.216
Local ridge	6.204	5.588	5.224	4.801	4.396
Temperature ridge	6.204	5.592	5.034	4.697	4.525
Residual ridge	6.217	5.580	5.207	5.375	4.727
Shrinkage blend	6.179	5.549	5.020	4.806	4.627

Additional OCM / ocm_f1 / random

Strategy	k=1	k=2	k=3	k=5	k=10
Global	6.120	6.120	6.120	6.120	6.120
Support mean	5.465	4.907	4.669	4.490	4.343
Offset	5.688	5.032	4.798	4.606	4.435
Affine	5.688	4.936	4.651	4.410	4.205
EB offset	5.310	4.964	4.787	4.621	4.451
Local ridge	5.465	4.492	4.022	3.593	3.252
Temperature ridge	5.465	4.742	4.106	3.779	3.580
Residual ridge	5.688	4.690	4.188	3.716	3.319
Shrinkage blend	5.428	4.469	3.996	3.552	3.230

Additional OCM / ocm_f1 / temperature_high

Strategy	k=1	k=2	k=3	k=5	k=10
Global	5.140	5.140	5.140	5.140	5.140
Support mean	6.204	5.886	5.712	5.556	5.436
Offset	6.300	5.816	5.588	5.397	5.255
Affine	6.300	5.760	5.527	5.328	5.208
EB offset	4.799	4.863	4.885	4.942	5.010
Local ridge	6.204	5.310	4.814	4.365	4.101
Temperature ridge	6.204	5.592	5.034	4.697	4.525
Residual ridge	6.300	5.307	4.861	4.454	4.233
Shrinkage blend	6.136	5.194	4.739	4.335	4.123

9 Mechanism and robustness interpretation

The external alignment diagnostic uses ten support draws, twenty permutations and five labels. Labels are permuted within the fixed support set; input records and their mean response remain unchanged. Mean, offset and EB offset therefore remain numerically invariant. The local F1 penalty is 1.2649 pp [0.8769,1.6941]; this ten-draw diagnostic is not substituted for the fifty-draw headline error. It diagnoses correspondence rather than proving chemical causality or computing a valid permutation-test p-value.

Same-alpha F1 residual differences are $-.2039$ $[-.4254, .0036]$ internally and .1230 $[-.0710, .3156]$ externally. The external alpha-matched point equals the primary residual point because both selected alphas are 1. F0 differences are .0095 internally and .0346 externally, with intervals spanning zero. The exact decomposition does not turn these inconclusive comparisons into positive evidence.

Applying all 25 old F1 outer-training fits to all additional sources at $k = 5$ gives random local/blend MAEs 3.6421/3.5797; difference $-.0624$ $[-.1603, .0244]$. High-temperature values are 4.5258/4.4357, difference $-.0901$ $[-.2066, .0167]$. These are error averages over fits, not prediction averaging, and do not replace the frozen full-fit primary result. Only three main sources report an unseen species and nine contain an operating value beyond an old input range. These coverage summaries are descriptive and do not validate a source selection rule.

10 Published-study weighting correction and reproduction scope

Archived Suvarna XGBoost predictions on 1,234 records and 109 sources yield record/source MAEs 57.8969/124.0426. The corresponding stacked raw-STY R^2 values are .8756/.3193. This release audits saved predictions; older full training-flow reproduction claims are not restated as newly executed results.

The Xu re-run uses 1,757 DOI-known rows from 156 sources; 82 DOI-missing upstream rows are outside the main audit. Its fixed XGBoost has 500 trees, depth 6, learning rate .05 and random state 42; the response is fitted in log1p space and transformed back for raw-unit MAE. Category maps and medians are fitted on each training partition. DOI is never a model input. Record and source five-fold fits use the same five seeds and current XGBoost 3.2.0. Pool record errors within each source/seed first; averaging sources directly from uneven source–fold cells produces a different weighting. Correct MAEs 69.6198/126.3649 replace the old five-fold grouping claim. Archived random-80/20 source-equal MAE is separately 62.9776. Old statements giving local OCM-F1 an average allocation advantage over affine are withdrawn; the present study's claims use its new source-tuned comparisons.

11 Reproduction files

Verified original-stage archives supply the analyses in Sections 1–11. The preceding acquisition-stage package independently contains all experiments in Section 12. The sixth release independently

contains all experiments in Section 13, including its inputs, predecessor settings and prediction exports. The current seventh release independently reproduces Section 14; earlier full releases remain separate, with exact identifiers and hashes. A final public deposit should combine these stages.

The frozen historical benchmark code is unchanged and hashes its inputs, splits and protocol. Repeated cells are preserved per fit with an index rather than duplicated in one large aggregate; the verified runner reconstructs an aggregate if requested. Historical row predictions are retained for support repeat zero, while every repeat's scalar cells remain present. All 480,500 external query prediction rows across F0/F1 include each support draw and budget. Export audit recomputes cell MAEs from those rows, checks support/query disjointness, shared representation ledgers and exact one-label identities. An atomic gzip writer is used for new large outputs; every distributed gzip stream and archive entry is validated.

The statistical engine is reproducible from the delivered cleaned snapshots. Original historical extraction, archived third-party model fitting and a prospective experimental campaign are distinct scopes. System dependency versions and commands are in the package entry-point guide. Document and plot builders consume the measured tables; no simulated performance values are used. Source attributions and third-party rights are retained separately from generated analysis code.

12 Acquisition extension: derivation, choices and complete factors

12.1 Fixed-correction risk identity

Let $G = V_T - H_A V_A$. Under the stated model, $q_A - f_T = -G\delta + H_A\epsilon$. Independence and zero means remove the cross term, yielding $\mathbb{E}\|q_A - f_T\|^2 = \operatorname{tr}(GBG^\top) + \sigma^2 \operatorname{tr}(H_A H_A^\top)$. A query-noise term, if independent and homoscedastic, adds a support-independent constant. The selector evaluates the latent-error expression. The 100,000-trial Gaussian simulation differs from the analytical risk by 0.293%. This verifies the formula under the model, not its calibration on catalyst data. Label-updated component weights make the predictor nonlinear in the support responses, so the same formula is only a design surrogate for that predictor.

The ridge map is $H_A = \mathbf{1}_T \mathbf{1}_A^\top/k + (Z_T - \mathbf{1}_T \bar{z}_A^\top) Z_{A,c}^\top (Z_{A,c} Z_{A,c}^\top + \alpha I)^{-1}$. It satisfies $H_A \mathbf{1}_A = \mathbf{1}_T$. The leave-one-out evidence has the form $D = \operatorname{diag}(1 - h_{ii})^{-1}(I - H_{AA})U_A$. Since $(I - H_{AA})\mathbf{1} = 0$, its rank is at most $k - 1$. Constants on supports have no direct independent coefficient evidence. Correlation encoded in a prior can still influence posterior predictions; this is not a claim of complete information absence in every Bayesian model.

12.2 Historical covariance construction

Only historical sources with at least 15 records enter the covariance proxy. Component scales use the source-equal first and second moments of historical cross-fitted response functions. For each source s, center $F_s = U_s / \operatorname{diag}(s_j)$ and y_s, then fit

$$\theta_s = (F_{s,c}^\top F_{s,c} + 0.1 n_s I)^{-1} F_{s,c}^\top y_{s,c}.$$

The residual proxy is $\sigma_s^2 = \|y_{s,c} - F_{s,c}\theta_s\|^2 / \max(n_s - 1 - d_s, 1)$ with ridge effective slope degrees of freedom d_s. A positive floor of 10^{-8} prevents numerical degeneracy. The within-source working posterior covariance is $V_s = \sigma_s^2 (F_{s,c}^\top F_{s,c} + 0.1 n_s I)^{-1}$. With S_θ the across-source sample covariance, set $B = 0.5 S_\theta + 0.5 \operatorname{diag}(S_\theta) + \operatorname{mean}_s V_s$ and $\sigma^2 = \operatorname{mean}_s \sigma_s^2$. These are preset empirical-Bayes proxies rather than unbiased estimates of physical noise or a proven optimal covariance.

For historical control tuning, direct covariance records from the pseudo-source's leakage group are removed. Historical cross-fitted dictionaries and inherited hyperparameters are formed within the entire outer historical-training set. This tuning stage is not a new independent test of the surrogate, and does not claim full repeated nesting of every inherited choice. The actual outer query groups and all additional-publication outcomes are excluded from all fitted covariance and control parameters.

12.3 Policies, nested budgets and query access

All selectors sort candidates by stable row identifier; exact ties choose the first. Coverage starts at the candidate nearest the pool feature centroid, then chooses a maximin point. Temperature coverage uses only standardized temperature; feature coverage uses all training-transformed features. Linear V-optimal design uses features $[10, z]$, a unit coefficient prior, and observation variance equal to the inherited local ridge α. Augmented V-optimal design uses $[10, z, U/s]$ and inherited ordinary augmented-ridge α. Its three component features are shared with the new response dictionary. Greedy V-optimal design minimizes mean query posterior variance. This is an ordinary standard-design implementation, not a claimed bitwise reproduction of Yu et al.'s complete algorithm or Fontaine et al.'s heteroscedastic online algorithm.

Innovation risk uses the exact local-ridge map in the main-text equation. Isotropic risk replaces B by $\operatorname{tr}(B)I/3$. No-component risk sets $B = 0$, leaving noise amplification only. No-noise risk sets $\sigma^2 = 0$. Pool-only risk replaces the query feature set by the complete pool. Prefix sizes $1, 2, 3, 5, 10$ are reported; intermediate sizes use the regularizer at the next reported budget. This rule maintains a common nested order. The sequence is greedy and need not minimize the final-budget global objective. V-optimal and the four query-aware risk rules see query covariates; coverage, random and pool-only risk do not. No rule sees acquisition-pool or query outcomes when selecting positions.

12.4 Predictors and historical retuning

The thirteen predictors are: global; support mean; frozen local ridge; frozen original blend; ANOVA transfer; ANOVA fixed prior; frozen local RBF; frozen ordinary two-component augmented ridge; same-local-alpha residual ridge; policy-retuned local ridge; policy-retuned blend prior; policy-retuned local RBF; and policy-retuned ordinary three-component augmented ridge. The inherited two-component features use empirical Shapley allocation with its own paired-reference intercept, rather than substituting the product-reference ANOVA functions. Their predictions remain unchanged under random acquisition to numerical precision.

Retuned local/ordinary augmented ridge choose $\alpha \in \{0.1, 1, 10, 100\}$. RBF chooses $\alpha \in \{0.01, 0.1, 1\}$ and $\gamma \in \{0.1, 1, 10\}$ with squared distances divided by the transformed feature dimension. Blend retuning chooses prior weight $0, 0.25, 0.5, 0.75, 1$, retaining its original two regularizers and shrinkage penalty. Tuning uses source-equal pseudo-query MAE, ten random orders and one order for each deterministic rule, entirely inside historical training groups. Settings are selected separately by policy and budget. All selected values and complete candidate scores are archived.

The ANOVA adapter uses inherited prior grids $\{0, 0.5, 1\}^3$ and penalties $\{\infty, 10, 2, 0.5, 0.1\}$ from the preceding response study. Its direction-aware precision is $\Omega = cD_0^2(I + D^\top D/D_0^2)^{-1}$, and box-constrained weights minimize $w^\top(D^\top D + \Omega)w - 2w^\top(D^\top e + \Omega\mu)$. Exact active-set enumeration is used in three dimensions. Infinite c returns the prior. This remains a retrospective component-transfer candidate, rather than a newly established estimator family.

12.5 Complete acquisition factors and verification

The extension contains 26 fitted regimes, 66,965 policy/budget error cells and 130 predictor–policy combinations. Nine inherited methods on random supports reproduce 510,750 error values to a maximum difference of 3.56×10^{-15}. Independent auditing re-aggregates 7,367,035 row-prediction values and re-computes 1,248,650 values directly from support labels. Every one of 13,393 support orders is reproduced after poisoning all evaluated outcomes, proving label-free acquisition for the frozen deployment interface. This does not erase retrospective corpus inspection.

Four candidate-focused random-query $k = 5$ contrasts form the reported adjusted family: innovation-risk ANOVA versus random ANOVA; versus augmented-design ANOVA; versus augmented-design retuned blend; and versus innovation-risk retuned blend. Percentile tails are $0.05/(2 \times 4)$ and its complement. Most other intervals remain nominal. Draws, seeds and folds are not independent sources. The covariance of training-fit uncertainty is not included in source-group intervals.

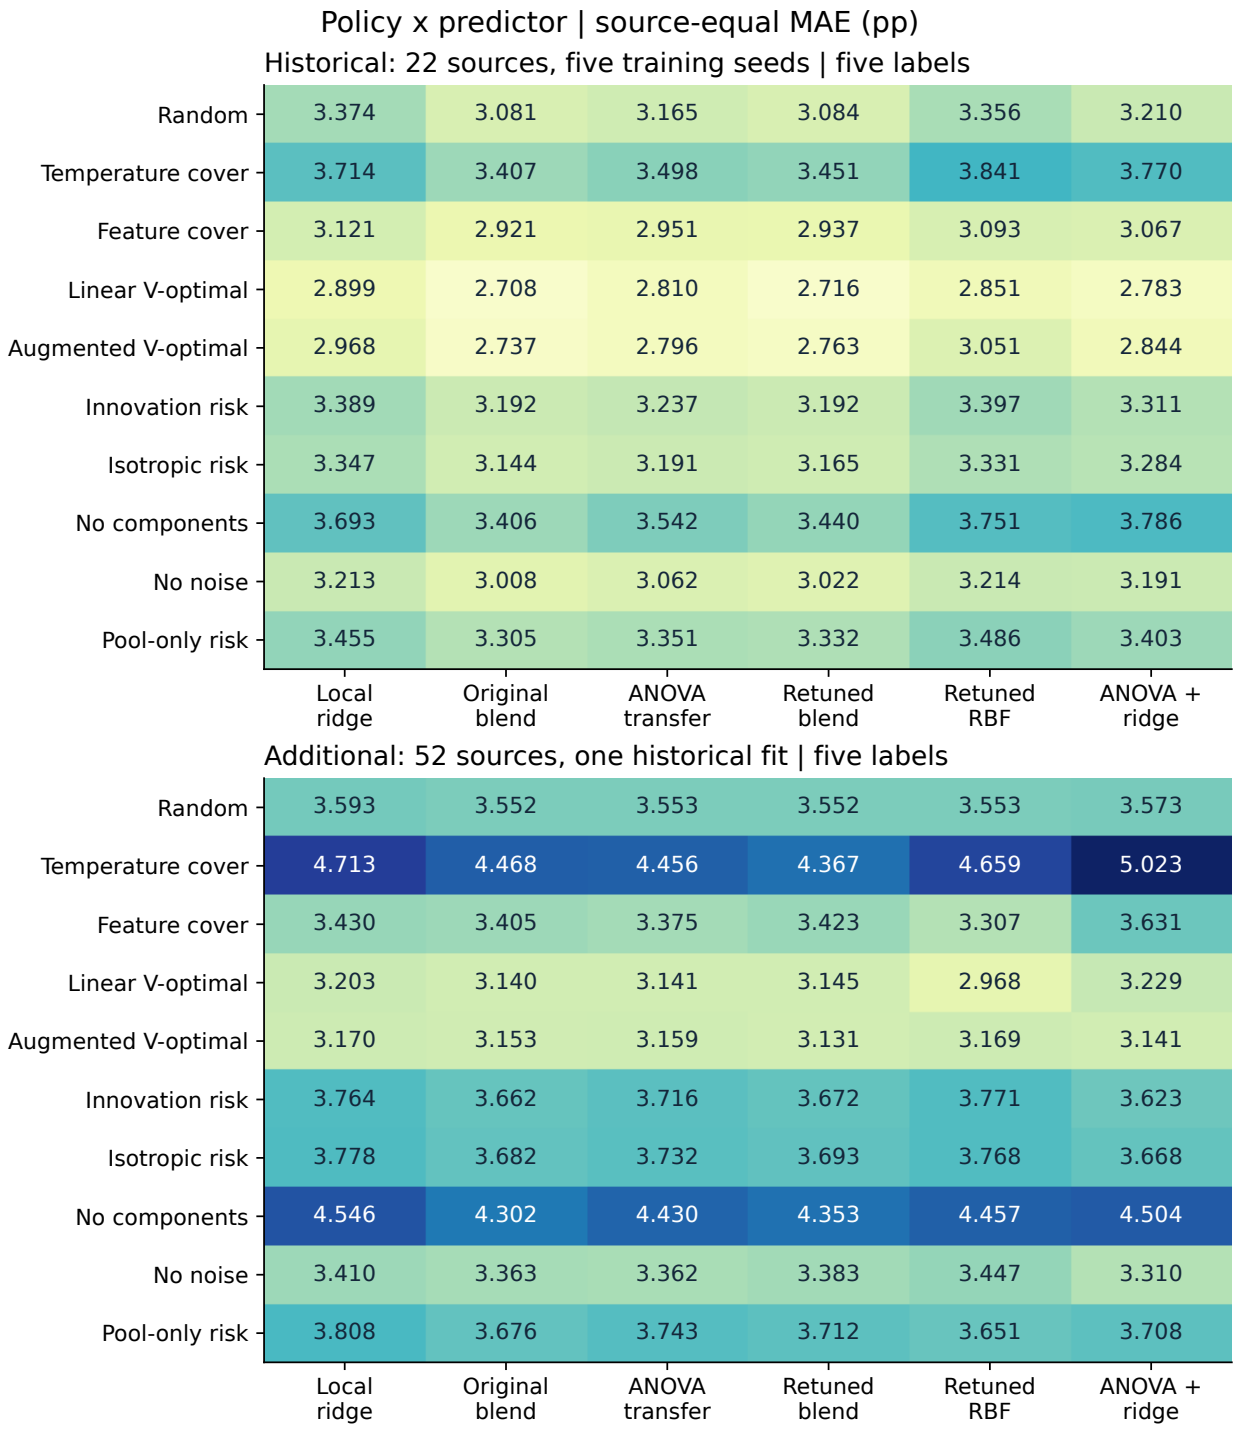


Figure 1: Selected columns of the complete acquisition factorial at five labels. The following tables retain all thirteen predictors and ten policies.

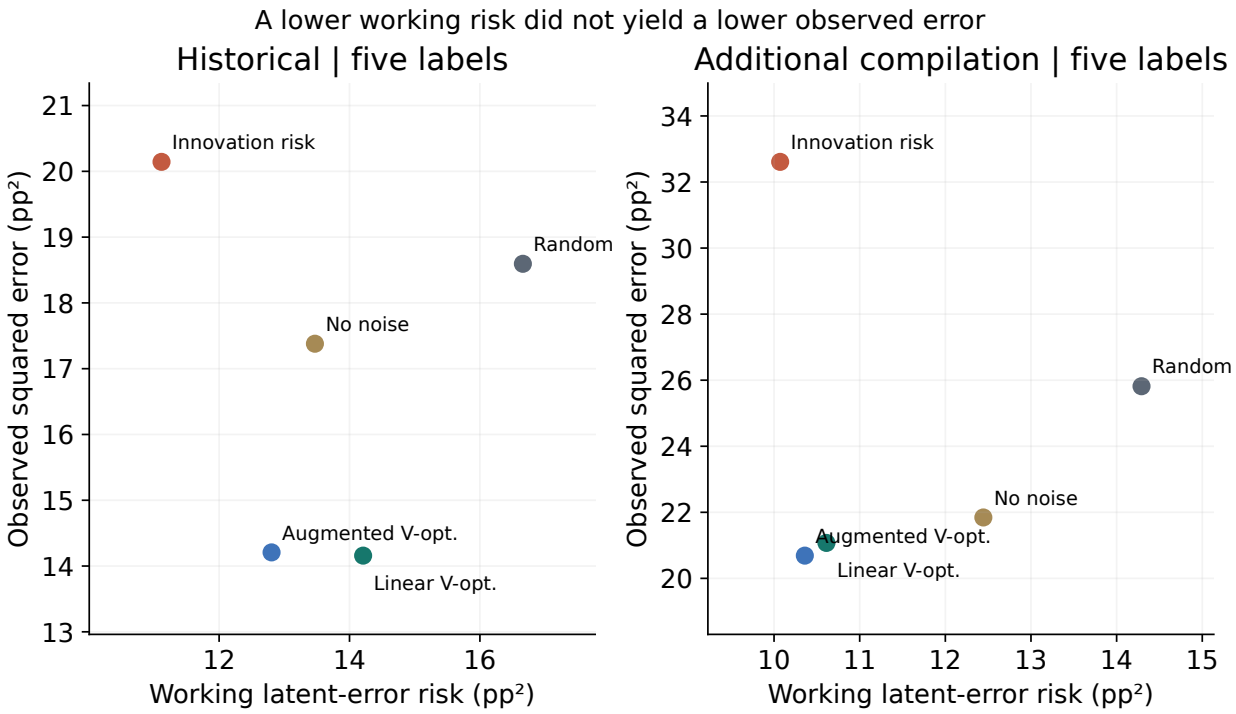


Figure 2: Secondary risk-calibration diagnostic for the fixed-prior predictor. Latent working risk excludes irreducible query noise; observed squared errors include outcome variability. Within a source, this support-independent noise cannot explain the full ordering discrepancy.

Table 3: Historical random queries, 22 sources. Five-label MAE (pp), all factors.

Predictor	R	T	F	L	A	I	S	NC	NN	P
Global	4.527	4.527	4.527	4.527	4.527	4.527	4.527	4.527	4.527	4.527
Support mean	3.693	3.935	3.491	3.409	3.497	3.672	3.666	3.947	3.709	3.775
Local ridge	3.374	3.714	3.121	2.899	2.968	3.389	3.347	3.693	3.213	3.455
Original blend	3.081	3.407	2.921	2.708	2.737	3.192	3.144	3.406	3.008	3.305
ANOVA transfer	3.165	3.498	2.951	2.810	2.796	3.237	3.191	3.542	3.062	3.351
ANOVA fixed prior	3.162	3.452	2.986	2.803	2.831	3.264	3.205	3.465	3.064	3.362
Local RBF	3.356	3.755	3.143	2.899	3.047	3.435	3.370	3.720	3.256	3.484
2-component augmented	3.127	3.521	3.028	2.718	2.736	3.239	3.190	3.524	3.067	3.313
Same-alpha residual	3.171	3.420	2.961	2.802	2.795	3.261	3.150	3.462	3.032	3.389
Retuned local	3.374	4.011	3.114	2.823	3.046	3.355	3.332	3.750	3.319	3.492
Retuned blend	3.084	3.451	2.937	2.716	2.763	3.192	3.165	3.440	3.022	3.332
Retuned RBF	3.356	3.841	3.093	2.851	3.051	3.397	3.331	3.751	3.214	3.486
3-component augmented	3.210	3.770	3.067	2.783	2.844	3.311	3.284	3.786	3.191	3.403

R: random; T: temperature cover; F: feature cover; L: linear V-optimal; A: augmented V-optimal; I: innovation risk; S: isotropic risk; NC: no components; NN: no noise; P: pool-only risk. Each column shares supports across predictors.

Table 4: Additional random queries, 52 sources. Five-label MAE (pp), all factors.

Predictor	R	T	F	L	A	I	S	NC	NN	P
Global	6.120	6.120	6.120	6.120	6.120	6.120	6.120	6.120	6.120	6.120
Support mean	4.490	5.665	4.444	4.106	4.200	4.279	4.154	5.122	4.434	4.644
Local ridge	3.593	4.713	3.430	3.203	3.170	3.764	3.778	4.546	3.410	3.808
Original blend	3.552	4.468	3.405	3.140	3.153	3.662	3.682	4.302	3.363	3.676
ANOVA transfer	3.553	4.456	3.375	3.141	3.159	3.716	3.732	4.430	3.362	3.743
ANOVA fixed prior	3.553	4.456	3.375	3.141	3.159	3.716	3.732	4.430	3.362	3.743
Local RBF	3.553	4.674	3.307	2.968	2.961	3.558	3.679	4.457	3.233	3.651
2-component augmented	3.556	4.752	3.463	3.129	3.113	3.650	3.684	4.513	3.321	3.671
Same-alpha residual	3.716	4.351	3.501	3.252	3.269	3.747	3.762	4.516	3.429	3.797
Retuned local	3.593	4.713	3.455	3.155	3.170	3.764	3.778	4.546	3.410	3.808
Retuned blend	3.552	4.367	3.423	3.145	3.131	3.672	3.693	4.353	3.383	3.712
Retuned RBF	3.553	4.659	3.307	2.968	3.169	3.771	3.768	4.457	3.447	3.651
3-component augmented	3.573	5.023	3.631	3.229	3.141	3.623	3.668	4.504	3.310	3.708

R: random; T: temperature cover; F: feature cover; L: linear V-optimal; A: augmented V-optimal; I: innovation risk; S: isotropic risk; NC: no components; NN: no noise; P: pool-only risk. Each column shares supports across predictors.

13 Source-contrast covariance transfer and matching acquisition

13.1 Why change the uncertainty model?

The earlier acquisition risk identity is mathematically correct under its stated working model, but its empirical ordering does not track observed errors well. This is motivation for a distinct candidate, not an identified diagnosis of which earlier assumption failed. The present model directly describes cross-fitted errors of γf_G after removing source levels. It combines existing random-coefficient regression, a flat intercept, functional-response features and a matching integrated-variance objective. None of empirical Bayes, universal kriging, cross-fitting or V-optimality alone is an originality claim.

13.2 Feature scaling and contrast records

For historical source s with at least 15 compatible rows, the unscaled feature vector contains the six preprocessed operating covariates and three ANOVA response functions. Each row's backbone, preprocessing and ANOVA reference set excludes that row's leakage group. Let $a_j = S^{-1}\sum_s \mathrm{mean}_i\, \phi_{sij}$ and $d_j^2 = S^{-1}\sum_s \mathrm{mean}_i\, \phi_{sij}^2 - a_j^2$, floored at 10^{-8}. Divide by d_j; an additional subtraction of a_j is unnecessary because the unrestricted source intercept makes joint feature translations irrelevant. Operating-only covariance uses the first six dimensions. All other variants use nine. These semantic coordinates allow three OOF dictionaries with different composition vocabularies to be combined without merging their full encoded feature vectors.

An outer historical fit has three first-level source-group folds. In tuning fold j, the whole held fold is absent from the covariance corpus. Within its complement, three new source-group folds supply OOF predictions for covariance rows. This requires nine secondary backbones per outer fit. Direct in-sample records from the same allowed complement form the in-sample ablation. The final deployment covariance uses the original three first-level OOF dictionaries on the complete historical training set. A fresh full historical fit reproduces the preceding full-fit geometry to numerical precision. The actual outer query groups, additional-publication outcomes and their labels never enter these fits.

13.3 Stabilized posterior-moment estimation

Define $X_s = \Phi_s - \mathbf{1}\bar{\phi}_s^\top$ and $t_s = (y_s - \gamma f_{G,s}) - \mathbf{1}\,\mathrm{mean}(y_s - \gamma f_{G,s})$. The working prior is zero-mean with covariance B. Given current B and σ^2,

$$V_s = (B^{-1} + X_s^\top X_s/\sigma^2)^{-1}, \tag{1}$$

$$m_s = V_s X_s^\top t_s/\sigma^2, \tag{2}$$

$$\widetilde{B} = S^{-1}\sum_s (V_s + m_s m_s^\top), \tag{3}$$

$$\sigma^2_{\mathrm{next}} = \frac{\sum_s[\|t_s - X_s m_s\|^2 + \mathrm{tr}(X_s^\top X_s V_s)]}{\sum_s (n_s - 1)}. \tag{4}$$

The full/operating/in-sample variants set $B_{\mathrm{next}} = 0.8\widetilde{B} + 0.2\,\mathrm{diag}(\widetilde{B})$; diagonal keeps only its diagonal, and isotropic sets $B_{\mathrm{next}} = \mathrm{tr}(\widetilde{B})I/p$. Eigendecomposition symmetrizes B and floors its eigenvalues. With $v = \max[S^{-1}\sum_s \mathrm{mean}(t_s^2), 10^{-6}]$, initialize $B = vI/p$ and $\sigma^2 = v/2$. The common positive floor is $\max(v \times 10^{-6}, 10^{-8})$. Stop after 80 updates or maximum relative covariance/noise change below 10^{-5}. Every iteration count and final change is released. The procedure is a shrunk posterior-moment fit, not a proven exact REML maximum-likelihood optimizer. Noise pools contrast degrees of freedom, whereas covariance averages posterior second moments per source. Source-equal fitting of every variance term is not claimed.

For each variant, nine pairs (γ, η) are assessed using ten random pseudo-support orders and budgets $1, 2, 3, 5, 10$. MAEs are first averaged over random orders, then equally over budgets and pseudo-sources. The minimum selects one pair for every budget and policy. The final covariance is refitted to allowed historical records at that gamma. This stage's pseudo-query scores are optimization

diagnostics. Backbone configuration and original standard-design regularizers are inherited, not independently nested again through a new global configuration search.

13.4 Established posterior and its information boundary

The main-text posterior follows the GP with an explicit constant basis and an infinitely diffuse intercept prior [17]. In a coefficient formulation, eliminating the intercept gives

$$V_A = \left[B^{-1} + \frac{\Phi_{A,c}^\top \Phi_{A,c}}{\nu}\right]^{-1}, \qquad m_A = V_A \Phi_{A,c}^\top (r_A - \overline{r}_A \mathbf{1})/\nu.$$

The predictor is $\gamma f_{G,Q} + \overline{r}_A + (\Phi_Q - \overline{\phi}_A) m_A$. For $B = \tau^2 I$ this is residual ridge with an unpenalized intercept and $\alpha = \nu/\tau^2$. Thus the isotropic version is an ordinary ridge working model; its mean transfer and its learned feature scaling are also explicit parts of the comparator.

The data information increment is $\Phi_{A,c}^\top \Phi_{A,c}/\nu$, with rank at most $\min(p, k-1)$. A component constant over supports has no independent contrast information. Off-diagonal prior covariance can couple its posterior update to observed directions, but those relationships originate in historical data, not new independent target labels. Bayesian conditioning cannot manufacture missing target information. At one label the posterior coefficient mean is zero and the prediction uses a global-response increment plus the estimated source offset.

With frozen working hyperparameters, additional support observations cannot increase latent posterior covariance in positive-semidefinite order; the flat-intercept limit is well defined for nonempty supports and positive noise. Consequently, the integrated-variance criterion is nonincreasing along the nested greedy sequence. Greedy selection need not be globally optimal, nor does decreasing working variance imply decreasing real MAE. The observation predictor's working variance adds ν to the latent diagonal. Neither variance includes uncertainty in estimated B, gamma, feature dictionary or inherited backbone choice.

13.5 Controls and statistical estimands

Policies are random, inherited linear V-optimal, inherited ANOVA-augmented V-optimal, contrast V-optimal, and contrast pool-only. Candidate pools and query rows are identical across rules. The two standard-design orders are exactly reproduced using the preceding regularizers and component scales. The contrast rule sees known query features; pool-only does not. Random is evaluated with 50 orders per source, deterministic rules once. All predictors share each policy's labels.

The eight noncandidate predictors are global, support mean, frozen local ridge, frozen original blend, re-tuned blend, re-tuned local ridge, re-tuned RBF, and ordinary ANOVA-feature ridge. Re-tuned ridge/residual controls use $\alpha \in \{0.1, 1, 10, 100\}$. RBF uses the preceding nine $(\alpha, \gamma_{\mathrm{RBF}})$ combinations. Ordinary ANOVA ridge uses inherited component scales and four alphas. Stronger blend controls first select separate local/residual alphas by source-equal pseudo-query MAE, then select a prior weight in $\{0, 0.25, 0.5, 0.75, 1\}$ and shrinkage in $\{\infty, 10, 2, 0.5, 0\}$; its evidence scale remains inherited. These controls are tuned separately per policy/budget, whereas every covariance variant keeps one random-tuned pair across all budgets/policies.

All evaluation MAEs first average orders within each source/seed, then seeds within source, then sources. Publication leakage groups are resampled in 20,000 bootstrap draws. Intervals condition on fitted models. The four locked random-query $k = 5$ comparisons are contrast-design/contrast-predictor versus linear-design/re-tuned-blend; versus augmented-design/re-tuned-blend; versus contrast-design/re-tuned-blend; and versus augmented-design/the same contrast predictor. Bonferroni percentile tails are $0.05/(2 \times 4)$ and their complement, adjusted separately within each cohort. All additional contrasts are exploratory. Prior inspection of the same corpora remains a limitation despite this within-stage protocol lock.

Support permutation is a descriptive mechanism diagnostic using repeat-zero five-label supports, twenty permutations per cell. It preserves the support mean but removes its response–input

assignment. It does not estimate prospective deployment error, identify causal catalysis mechanisms or add independent sources.

13.6 Observed performance, matched ablations and verification

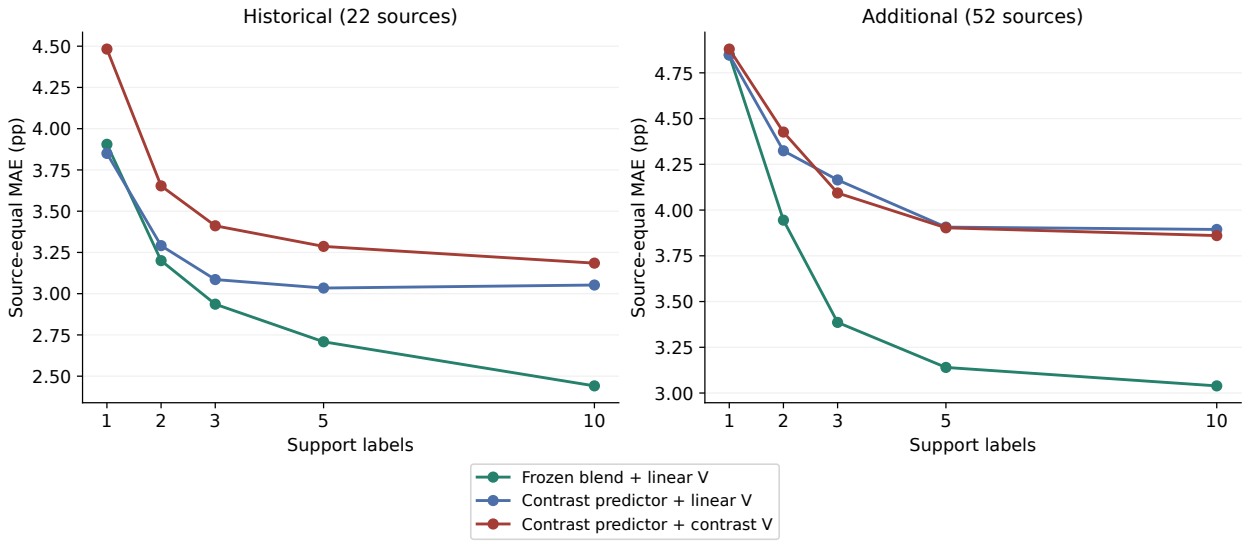


Figure 3: Random-query budgets with fixed source populations. Standard-design frozen blending remains a comparator at every budget; no curve is selected using additional-source outcomes.

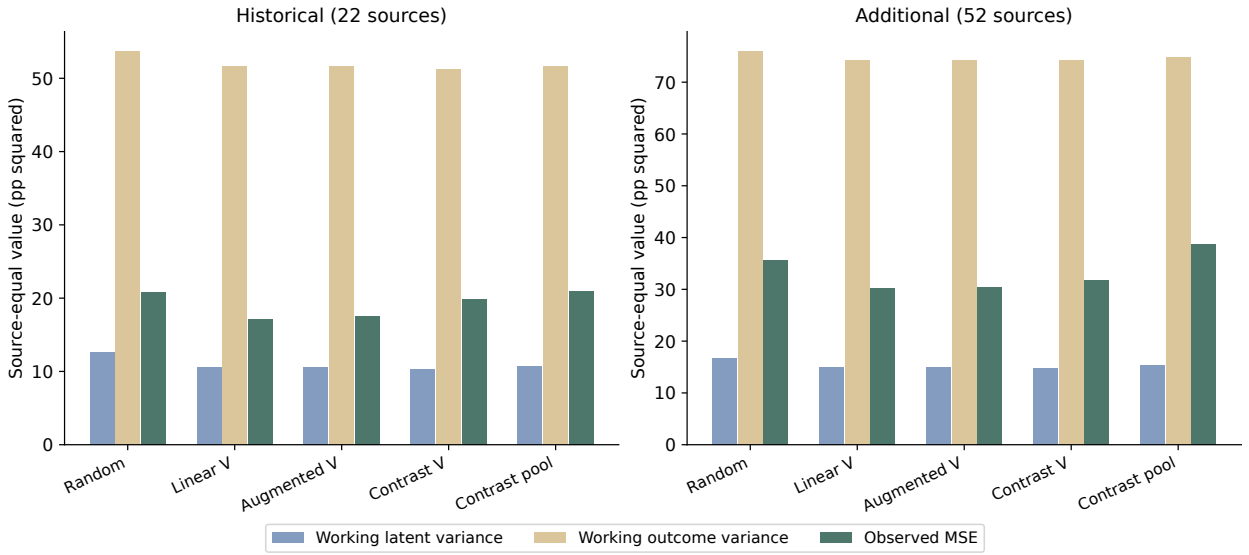


Figure 4: Five-label working variances versus observed MSE. Observation variance adds the frozen working residual term; this term includes lack of fit. These averages do not certify predictive coverage or identify physical laboratory noise.

Table 5: Historical, 22 sources. Five-label random-query MAE (pp), all 65 factors.

Predictor	R	L	A	C	P
Global	4.527	4.527	4.527	4.527	4.527
Support mean	3.693	3.409	3.497	3.624	3.675
Frozen local ridge	3.374	2.899	2.968	3.196	3.300
Frozen blend	3.081	2.708	2.737	2.920	3.030
Retuned blend	3.073	2.696	2.897	3.027	3.080
Retuned local ridge	3.374	2.823	3.046	3.304	3.297
Retuned RBF	3.356	2.851	3.051	3.303	3.222
Ordinary ANOVA ridge	3.210	2.783	2.844	3.129	3.230
Contrast full	3.365	3.034	3.090	3.287	3.389
Contrast diagonal	3.368	3.049	3.096	3.296	3.390
Operating only	3.430	3.064	3.153	3.340	3.452
Isotropic	3.390	3.059	3.107	3.300	3.413
In-sample covariance	3.382	3.065	3.063	3.291	3.369

R: random; L: linear V-optimal; A: augmented V-optimal; C: contrast V-optimal; P: contrast pool-only. All columns share query populations; each column shares supports across predictors.

Table 6: Additional, 52 sources. Five-label random-query MAE (pp), all 65 factors.

Predictor	R	L	A	C	P
Global	6.120	6.120	6.120	6.120	6.120
Support mean	4.490	4.106	4.200	4.015	4.290
Frozen local ridge	3.593	3.203	3.170	3.447	3.744
Frozen blend	3.552	3.140	3.153	3.411	3.710
Retuned blend	3.571	3.149	3.157	3.422	3.619
Retuned local ridge	3.593	3.155	3.170	3.447	3.744
Retuned RBF	3.553	2.968	3.169	3.478	3.744
Ordinary ANOVA ridge	3.573	3.229	3.141	3.349	3.709
Contrast full	4.298	3.907	4.003	3.903	4.184
Contrast diagonal	4.338	3.952	4.050	3.920	4.201
Operating only	4.405	4.032	4.114	3.981	4.280
Isotropic	3.936	3.589	3.617	3.753	3.995
In-sample covariance	4.231	3.802	3.833	3.927	4.206

R: random; L: linear V-optimal; A: augmented V-optimal; C: contrast V-optimal; P: contrast pool-only. All columns share query populations; each column shares supports across predictors.

Table 7: Post-hoc parameter-matched covariance diagnostic. Linear-design supports at five labels; gamma and total observation noise are fixed to the main candidate. Nominal intervals are exploratory.

Covariance	Historical MAE	Additional MAE
Contrast full	3.034	3.907
Contrast diagonal	3.040	3.951
Operating only	3.054	4.027
Isotropic	3.048	3.823
In-sample covariance	3.045	3.982

The candidate's historical dictionary has nine dimensions, whereas full-feature local ridge/RBF uses six operating features plus historical composition vocabulary channels. The covariance ablations are cleaner tests of covariance choices than the comparison with those wider dictionaries; neither supplies a unique causal explanation for failure. Independently tuned variants may choose different gammas/noise multipliers. The matched diagnostic was added after the locked results were observed, uses the main historical gamma and total noise, and makes no new selection from query outcomes. It evaluates all five covariance alternatives, including unfavorable ones.

The complete stage contains 61290 error cells, 65 predictor–policy combinations, and five new covariance-transfer variants. All 11,804 inherited support orders and 236,080 original-method MAEs reproduce the preceding stage to maximum 3.56×10^{-15}. The covariance stage refits 338 backbones in total: 26 full, 78 first-level and 234 secondary. The released numerical verifier checks isotropic-ridge equivalence, feature/response shift identities, all ten greedy steps against scalar scores, constant-feature nonidentifiability, and a 100,000-trial Gaussian risk simulation (0.311% relative error). Independent row audits and poisoned-outcome support-order checks are archived with exact counts. Prediction files are lossless float64 NPZ, without object arrays or quantization.

The sixth release's covariance entry point re-fits Section 13; it does not claim to newly re-fit all earlier multi-task studies. Full covariance records and arrays remain in that sixth release. The seventh release preserves the covariance summaries and figures but re-fits the separate cross-task acquisition experiment in Section 14. A separate fresh-directory smoke check re-fits one historical regime from an extracted package; the full mode re-fits all 26. Manifest verification and numerical audit use released outputs without treating a completed cache as a new training run.

14 Cross-task standard acquisition extension

14.1 Scope, chronology and fixed populations

This seventh extension evaluates five panels from four outcomes in two reaction systems, not five independent reaction datasets. Conversion and selectivity share the same paired records; twenty publication identifiers overlap the STY collection; OCM F0/F1 use the same records. These dependencies remain relevant when interpreting breadth. The current protocol is timestamped and hashed before current outer evaluation, but all input corpora were already inspected in earlier development. All new results are retrospective. Only the five-label frozen-blend query-V-minus-random comparison across the five random-query panels is primary. The parameter-matched OCM geometry crossing was added after primary results and is explicitly exploratory.

Every outer seed/fold and backbone configuration is inherited verbatim from the preceding full-study benchmark. Current outer models are re-fitted and compared with the original zero-label predictions. Historical source-disjoint OOF responses are reused for policy-specific control tuning. Inner preprocessors are reconstructed using the same three historical training-group partitions; no pseudo-validation group fits its own feature preprocessor. There is no new inner backbone fit or new global-configuration search in this extension. OOF model configuration was selected within the complete outer historical-training set in the earlier benchmark; historical control scores remain optimization diagnostics, not independent test results. Outer query groups are excluded from every such choice.

Random queries, high-temperature queries and reported-stratum holdouts use the original label-blind definitions. Eligibility is unchanged: at least ten pool records, at least five query records. Pool IDs are sorted before selection. Deterministic acquisitions are evaluated once, rather than copied fifty times as artificial replicates. Random orders run fifty times. Every policy produces one nested sequence to budget ten; intermediate steps use the original local-ridge alpha of the next reported budget. Different target outcomes can have different frozen alpha schedules, even when they share feature rows.

14.2 Five acquisitions and thirteen predictors

Temperature coverage uses the first numeric feature; full-feature coverage uses all allowed standardized/imputed numeric and encoded-composition features. The first row is nearest the pool mean and later rows maximize their minimum squared distance to the selected set. This is a conventional coverage heuristic. Query and pool V-optimal acquisition use a unit linear coefficient prior on $[10, z]$ and the inherited random-design local alpha. Query V minimizes mean working posterior variance on query covariates; pool V minimizes it over all candidate-pool covariates, including candidates that may later be observed. Neither reads outcomes. The finite intercept prior used for acquisition differs from the exact unpenalized intercept used for prediction, so working variance is only a design surrogate.

Nine original predictors are frozen: global, support mean, residual offset, affine, empirical Bayes, local ridge, temperature ridge, residual ridge and shrinkage blend. Four stronger controls re-tune local ridge, residual ridge, blending and local RBF by policy/budget. Ridge alphas are $\{0.1, 1, 10, 100\}$. RBF alpha and gamma both use $\{0.01, 0.1, 1\}$ and $\{0.1, 1, 10\}$, respectively; squared distances divide by feature dimension. Kernel ridge centers its support kernel and response and restores the support mean. After local/residual alpha selection, blending evaluates five prior weights times five shrinkages, using the inherited response scale and support-only analytic LOO. Each candidate's MAE is averaged over historical support draws, then pseudo-sources. Random pseudo-supports repeat ten times; deterministic policies once. All control grids and acquired rows are common across policies. An ordinary kernel/local fit is a strong comparator; a lower point estimate is not selected using outer query outcomes.

14.3 Primary family, units and aggregation

Each error cell averages query rows, then repeated support draws are averaged within source/seed, then seeds within source, then sources receive equal weight. A leakage group may contain more than one source; cluster bootstrap retains those source weights. Twenty thousand group resamples produce nominal 95% intervals. Five-panel primary tails are 0.005 and 0.995, the approximate Bonferroni-adjusted 99% percentile interval. Adjustment accommodates endpoint dependence but does not include training-fit uncertainty or remove retrospective selection. No folds, random draws, or shuffled-label draws are counted as independent publications. STY uses its native productivity unit, while the other outcomes use percentage points; raw errors are not pooled across these units.

Table 8: Five-panel adjusted primary results. Same interpretation and units as the main primary table.

Panel	n	Random	Query V	Gain (%)	Adjusted Δ interval
Conversion	32	2.441	2.121	13.1	[-0.707, 0.068]
Selectivity	32	11.584	10.683	7.8	[-2.140, 0.615]
STY	28	62.608	50.274	19.7	[-25.965, -2.874]
OCM F0	22	3.660	3.768	-3.0	[-0.355, 0.714]
OCM F1	22	3.081	2.708	12.1	[-0.746, -0.056]

Table 9: Frozen blend across five acquisitions at five labels. R: random; T: temperature coverage; F: feature coverage; V: query V; P: pool V. STY has native productivity units; other rows are percentage points.

Panel	R	T	F	V	P
Conversion	2.441	2.647	2.172	2.121	2.417
Selectivity	11.584	13.606	11.096	10.683	10.481
STY	62.608	70.661	53.111	50.274	58.440
OCM F0	3.660	3.901	4.008	3.768	3.895
OCM F1	3.081	3.407	2.921	2.708	2.886

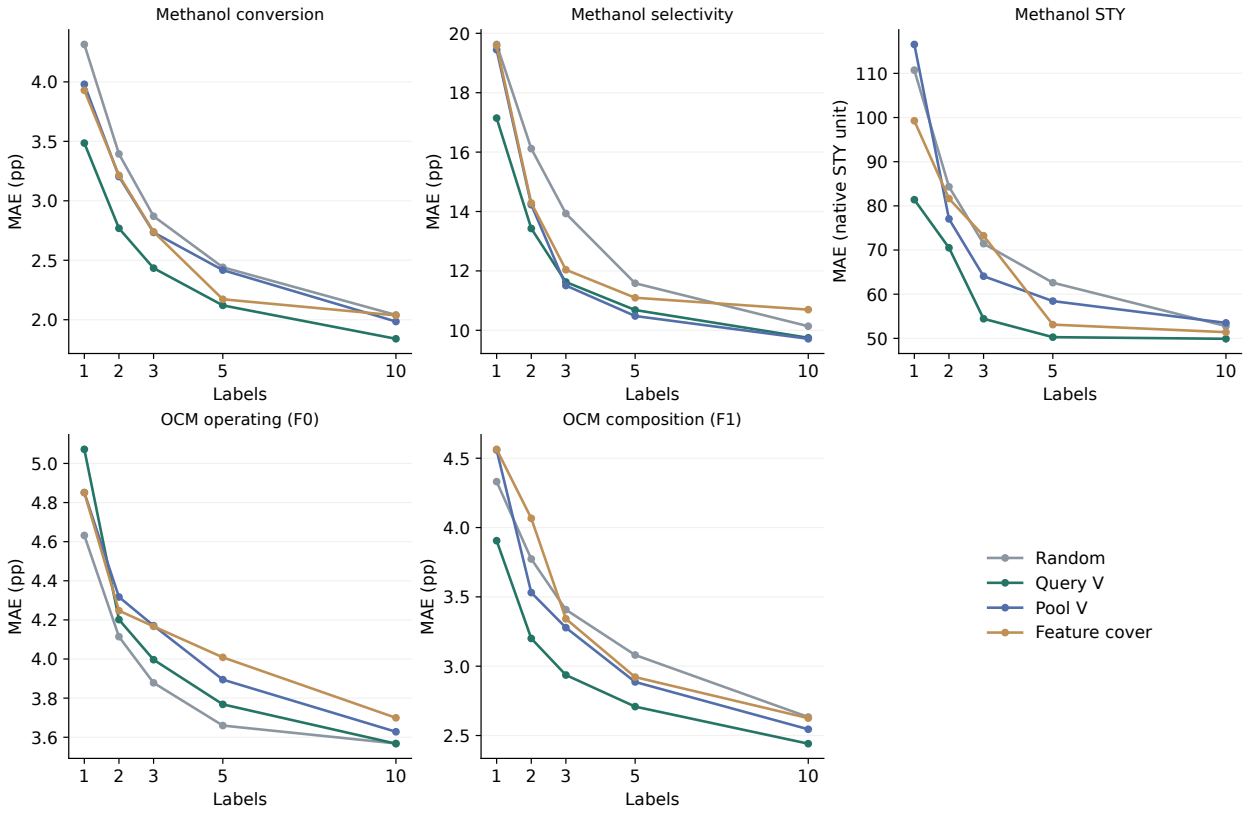


Figure 5: Budget curves for the unchanged frozen blend on five random-query panels. All rules share each panel's query population. The four displayed rules are representative; all five acquisitions and thirteen predictors are available in the numerical tables.

Table 10: Methanol conversion: all five-label random-query factors. Native STY units for STY; percentage points otherwise.

Predictor	R	T	F	V	P
Global	5.271	5.271	5.271	5.271	5.271
Support mean	4.304	4.182	4.063	3.806	4.315
Offset	3.455	3.459	3.376	3.226	3.539
Affine	3.083	3.239	2.926	2.839	3.098
Empirical Bayes	3.457	3.437	3.383	3.242	3.506
Frozen local ridge	2.429	2.576	2.197	2.135	2.394
Temperature ridge	2.519	2.711	2.470	2.343	2.587
Frozen residual ridge	2.720	2.873	2.540	2.395	2.669
Frozen blend	2.441	2.647	2.172	2.121	2.417
Retuned local ridge	2.429	2.647	2.271	2.173	2.472
Retuned residual ridge	2.720	2.856	2.524	2.421	2.668
Retuned blend	2.421	2.633	2.187	2.120	2.407
Retuned local RBF	2.401	2.531	2.075	2.153	2.436

R: random; T: temperature cover; F: feature cover; V: query V-optimal; P: pool V-optimal. Supports are common across predictors within a policy.

Table 11: Methanol selectivity: all five-label random-query factors. Native STY units for STY; percentage points otherwise.

Predictor	R	T	F	V	P
Global	22.740	22.740	22.740	22.740	22.740
Support mean	17.136	16.999	16.998	15.806	17.521
Offset	15.747	16.207	16.310	15.126	16.275
Affine	12.539	12.520	12.214	11.842	12.014
Empirical Bayes	15.988	16.439	16.481	15.424	16.501
Frozen local ridge	11.547	12.689	11.142	10.725	10.568
Temperature ridge	11.535	11.561	11.626	10.973	11.299
Frozen residual ridge	12.077	14.678	11.332	10.735	10.850
Frozen blend	11.584	13.606	11.096	10.683	10.481
Retuned local ridge	11.547	12.381	11.336	10.762	10.568
Retuned residual ridge	12.077	13.917	11.102	10.790	10.793
Retuned blend	11.623	12.159	11.097	10.691	10.579
Retuned local RBF	11.122	11.201	10.774	10.558	10.231

R: random; T: temperature cover; F: feature cover; V: query V-optimal; P: pool V-optimal. Supports are common across predictors within a policy.

Table 12: Methanol STY: all five-label random-query factors. Native STY units for STY; percentage points otherwise.

Predictor	R	T	F	V	P
Global	139.258	139.258	139.258	139.258	139.258
Support mean	91.993	110.199	101.256	77.514	97.591
Offset	88.265	104.855	92.493	78.964	90.901
Affine	71.225	82.290	69.447	62.853	71.968
Empirical Bayes	86.510	98.390	90.549	78.972	88.594
Frozen local ridge	65.210	72.805	56.486	52.671	61.580
Temperature ridge	87.098	96.672	86.843	66.248	88.206
Frozen residual ridge	69.889	76.255	62.666	57.442	65.811
Frozen blend	62.608	70.661	53.111	50.274	58.440
Retuned local ridge	65.210	69.453	56.283	51.293	60.031
Retuned residual ridge	69.889	77.153	62.784	57.352	65.539
Retuned blend	62.613	69.483	53.103	49.852	56.767
Retuned local RBF	63.957	68.185	54.535	51.547	62.927

R: random; T: temperature cover; F: feature cover; V: query V-optimal; P: pool V-optimal. Supports are common across predictors within a policy.

Table 13: OCM operating (F0): all five-label random-query factors. Native STY units for STY; percentage points otherwise.

Predictor	R	T	F	V	P
Global	5.302	5.302	5.302	5.302	5.302
Support mean	3.693	3.935	4.038	3.802	3.907
Offset	3.672	3.931	3.997	3.819	3.932
Affine	3.694	3.933	4.005	3.796	3.941
Empirical Bayes	3.759	3.963	4.000	3.866	3.956
Frozen local ridge	3.655	3.905	4.027	3.754	3.877
Temperature ridge	3.677	3.961	4.048	3.771	3.885
Frozen residual ridge	3.701	3.930	4.012	3.799	3.948
Frozen blend	3.660	3.901	4.008	3.768	3.895
Retuned local ridge	3.655	4.006	4.111	3.776	3.908
Retuned residual ridge	3.701	3.974	4.070	3.792	4.007
Retuned blend	3.665	3.989	4.072	3.774	3.945
Retuned local RBF	3.671	3.972	4.086	3.756	3.931

R: random; T: temperature cover; F: feature cover; V: query V-optimal; P: pool V-optimal. Supports are common across predictors within a policy.

Table 14: OCM composition (F1): all five-label random-query factors. Native STY units for STY; percentage points otherwise.

Predictor	R	T	F	V	P
Global	4.527	4.527	4.527	4.527	4.527
Support mean	3.693	3.935	3.491	3.409	3.688
Offset	3.423	3.682	3.195	3.041	3.429
Affine	3.360	3.564	3.193	2.976	3.365
Empirical Bayes	3.415	3.645	3.227	3.120	3.366
Frozen local ridge	3.374	3.714	3.121	2.899	3.109
Temperature ridge	3.677	3.961	3.471	3.373	3.664
Frozen residual ridge	3.109	3.332	2.940	2.759	2.927
Frozen blend	3.081	3.407	2.921	2.708	2.886
Retuned local ridge	3.374	4.011	3.114	2.823	3.114
Retuned residual ridge	3.109	3.554	3.056	2.790	2.961
Retuned blend	3.073	3.659	2.994	2.696	2.911
Retuned local RBF	3.356	3.841	3.093	2.851	3.116

R: random; T: temperature cover; F: feature cover; V: query V-optimal; P: pool V-optimal. Supports are common across predictors within a policy.

14.4 Label-correspondence diagnostic

At budget five, the first random order and each deterministic support are used for twenty outcome permutations per source/scenario/seed. The input rows, global predictions, support outcomes as a multiset, and exact support mean remain the same; only input–label correspondence changes. The frozen blend is re-computed, including its support-only blend weight. Query labels enter scoring only. These are exploratory controls, not repeated training datasets or a physical perturbation experiment. Permutation errors average shuffles and seeds before source-group intervals.

Table 15: Correspondence shuffling under query V: frozen blend, five labels. Nominal source-group intervals are exploratory.

Panel	Aligned	Shuffled	Shuffled minus aligned, 95% CI
Conversion	2.121	4.135	2.015 [1.253, 2.851]
Selectivity	10.683	16.723	6.040 [3.693, 8.513]
STY	50.274	80.945	30.671 [10.436, 59.603]
OCM F0	3.768	3.815	0.047 [-0.009, 0.134]
OCM F1	2.708	3.548	0.840 [0.478, 1.237]

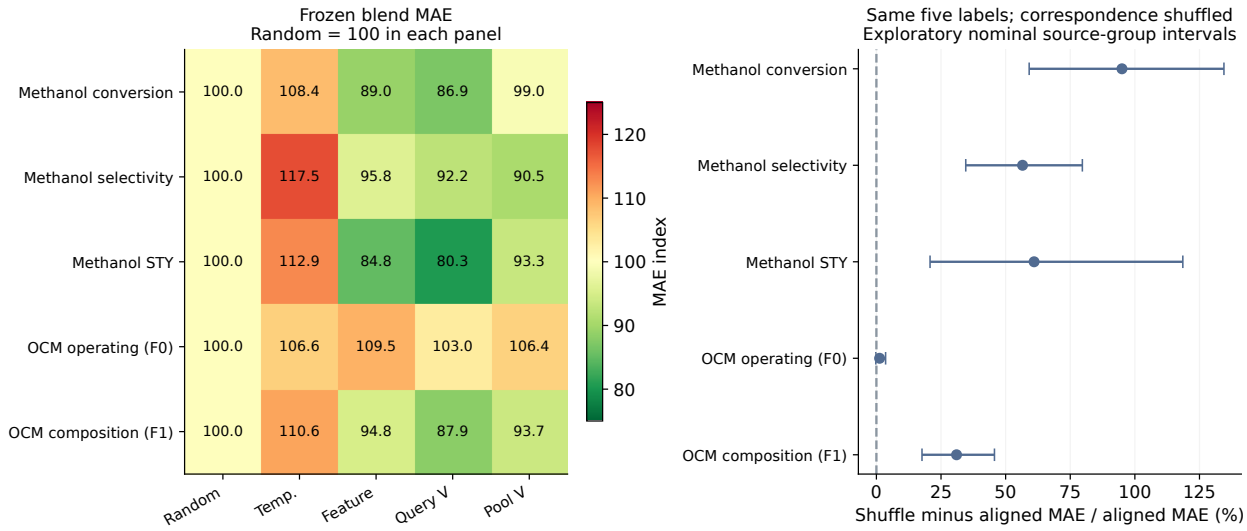


Figure 6: Left: all acquisition means for the frozen blend, normalized to each panel's observed random-design MAE. Right: outcome-correspondence shuffling effect, nominal native-unit intervals rescaled by the fixed observed aligned MAE. Normalization is for display, not a pooled or bootstrap ratio estimand.

14.5 Post-primary representation crossing

The OCM follow-up fixes the F1 alpha schedule for both operating-only (six-feature) and full composition-enriched V geometry. It evaluates each acquired sequence with both already frozen F0 and F1 predictors. Numeric preprocessors are verified equal between paired representations; outcomes, pool rows and queries are identical. The full-geometry orders exactly match the original F1 orders. The operating-only orders in this follow-up can differ from the primary F0 orders because the regularizer is now fixed to F1, rather than independently inherited from F0. A direct comparison of the primary F0 and F1 numbers would mix representation, backbone/adapter choices and acquisition regularization; this crossing removes the acquisition-alpha difference for its own conditional contrasts.

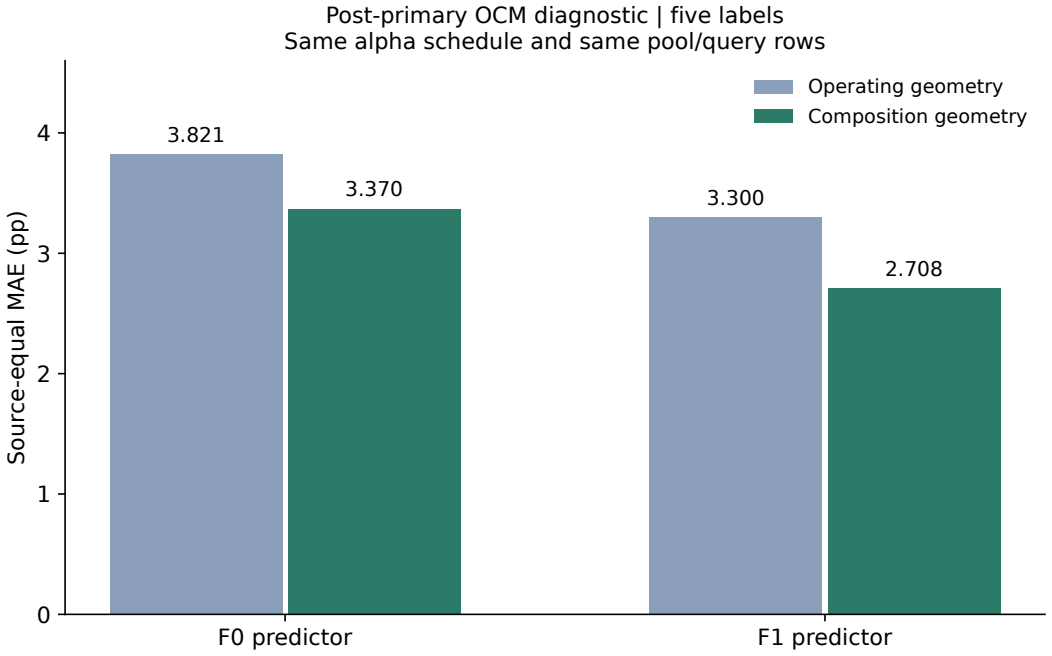


Figure 7: Post-primary exploratory crossing of selection geometry and prediction representation. Same F1 alpha, five supports, 22 OCM sources. Geometry-specific nominal intervals and all budgets are released. No new training or parameter selection occurs. Composition used for acquisition is not automatically an input to the F0 predictor.

14.6 Independent verification and reproduction scope

The current computation contains 125 outer fits, 330,750 error cells, 66,150 nested orders, and 31,712,850 scalar row predictions across 65 combinations per panel. Exactly 125 new outer backbones and zero new inner backbones are trained. Original random-design MAEs reproduce across 306,250 cells and nine methods, with maximum discrepancy 2.77e-09. The larger STY scale explains why this roundoff is not identical to the percentage-point panels. OCM frozen query/random orders and original four-method errors independently match the sixth release.

Independent audits re-aggregate 31,712,850 scalar predictions and reconstruct 2,936,375 from support rows. All 66,150 selection orders remain unchanged after replacing every target response by values between -10^9 and 10^9. Five hundred outer/inner group-separation checks pass. Mathematical verification compares all twenty greedy steps of query/pool design with a direct primal posterior and checks exact ridge prediction/LOO against independent scikit-learn fits. The post-primary crossing adds 2,700 cells and 172,800 scalar predictions, with an independent read-only audit.

The seventh release includes lossless float64 NPZ predictions, acquired/query ID ledgers, historical control scores, every parameter, input snapshots, split ledgers, predecessor references and a checksummed manifest. Its verify mode reads the release and writes audit reports only into a temporary directory. Smoke mode re-fits one OCM F1 outer fold in an empty external directory; full mode re-fits all 125 current outer models, re-tunes the current controls from inherited historical OOF records and repeats the post-primary geometry diagnostic. A completed cache is never counted as a fresh refit. The smoke check is a portability check, not a second full 125-fit independent repetition. Earlier Section 1–13 studies retain their complete raw arrays in the prior numbered releases, rather than being claimed as newly re-run by the current entry point.

15 Executed acquisition-view follow-up

15.1 Scientific scope and finite choices

This follow-up contains eleven acquisition policies and twelve predictors for each OCM prediction representation. The policies are random; operating V; composition V; feature coverage; pool V; kernel V; robust multi-view V; fixed historical best view; a zero-margin view gate; the empirical-margin conservative gate; and composition-correspondence-shuffled V. Nested budgets are 1, 2, 3, 5 and 10. Random acquisition has fifty target draws and ten historical tuning draws; chemistry shuffling has ten draws. The same support order is crossed with both F0 and F1. Gate development targets F0's frozen blend at five labels, and therefore does not claim to optimize every other predictor or budget.

The original nine response methods are joined by policy-retuned local ridge, blend and local RBF. Ridge penalties are $\{0.1, 1, 10, 100\}$. The local RBF uses the inherited nine (α, γ) combinations from $\alpha \in \{0.01, 0.1, 1\}$ and $\gamma \in \{0.1, 1, 10\}$. Blending uses the inherited finite prior-weight and penalty grid. Scores are averaged over draws and equally over historical pseudo-sources. The acquisition kernel's gamma is selected using historical F0 blend losses. The robust design is a finite greedy maximum-remaining-variance-fraction control; it is not globally optimized over all support subsets.

The operating kernel is $K_O(x, x') = 100 + x_O^T x'_O$, and composition V uses the corresponding full-feature linear kernel. Kernel V uses $K_C(x, x') = 100 + \exp(-\gamma \|x_C - x'_C\|^2 / d_C)$. For a support set A, the mean query working variance is

$$V(A) = \overline{K(q, q)} - \frac{1}{n_Q} \mathrm{tr}\{K_{QA}(K_{AA} + \alpha I)^{-1} K_{AQ}\}.$$

Robust selection minimizes $\max\{V_O(A)/V_O(\varnothing), V_C(A)/V_C(\varnothing)\}$ at each greedy step. The alpha schedule is inherited from F1 and shared by the two views. These posterior calculations are working design criteria, without a calibrated-risk assertion.

15.2 Empirical view gate and a bounded theoretical interpretation

The gate fits eight fixed response-free contexts: log pool size, log query size, log operating variation, log chemistry variation, log query chemistry novelty relative to the pool, chemistry variation unexplained by a pool-only ridge on operating features, log chemistry mean shift, and the Jaccard disagreement between the two five-point support sets. The ridge penalty is fixed at 10. Fewer than six historical groups trigger composition fallback. Otherwise group-held-out mean absolute meta errors supply the 80th-percentile margin q. Meta fitting assigns inverse within-group frequency weights. Original adapter settings had been selected on the entire outer-training set, so this margin is not fully nested calibration of every dependency.

Write $\delta = L_C - L_O$, with estimated difference $\widehat{\delta}$. A zero-margin sign choice has pointwise regret at most $|\widehat{\delta} - \delta|$ relative to the better of the two realized losses. For the empirical-margin rule O iff $\widehat{\delta} > q$, with $q \geq 0$, regret is at most $|\widehat{\delta} - \delta| + q$; the excess loss from an incorrect switch relative to staying with C is at most $\max\{|\widehat{\delta} - \delta| - q, 0\}$. These elementary identities follow directly by considering the two signs of δ. They do not control the meta prediction error under source shift, provide a coverage guarantee, or constitute a novel learning theorem. Transductive experimental design and kernel integrated-variance selection have substantial prior literature[13, 23]; the current evaluated contribution is the catalysis representation crossing and measured information decomposition.

Table 16: Locked historical primary contrasts for the conservative view gate.

Fixed control	Δ MAE	99% lower	99% upper
Composition V	0.0035	0.0000	0.0138
Operating V	-0.4484	-1.1429	0.0958
Kernel V	-0.0850	-0.2685	0.0419
Robust multi-view V	-0.0063	-0.0371	0.0147
Fixed historical best	0.0035	0.0000	0.0138

Candidate minus control; negative favors the gate. Twenty-two equally weighted sources. No contrast supports gate superiority after the five-comparison adjustment. Fixed historical best and composition V coincide here.

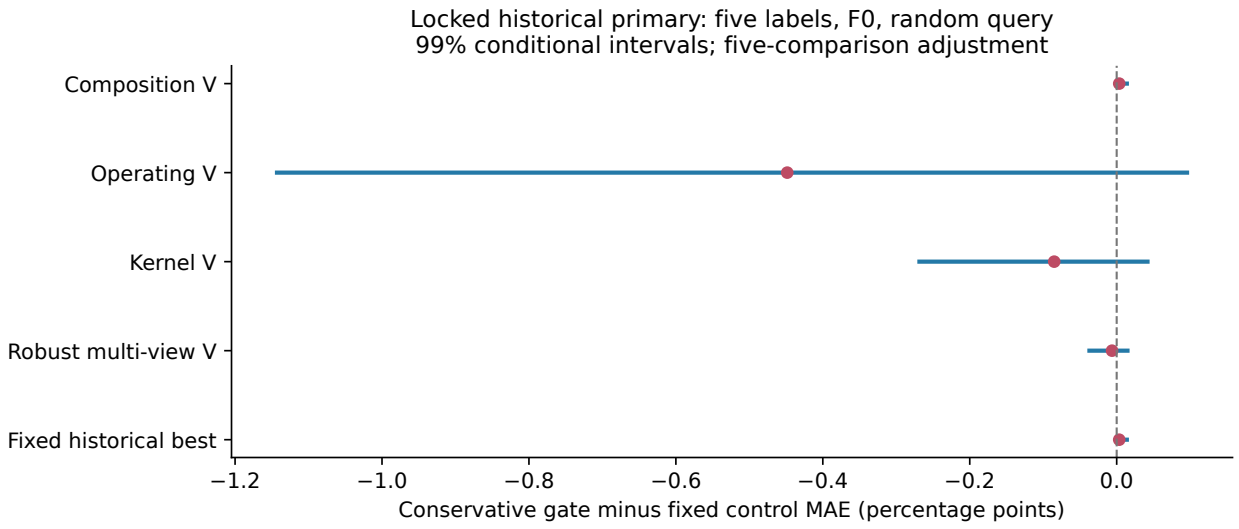


Figure 8: Historical primary gate-minus-control contrasts. Intervals are 99% conditional leakage-group bootstrap intervals for five-comparison adjustment; they average support draws and seeds before source-level resampling. No contrast supports superiority. The fixed historical best is composition V in every relevant fold, so those two controls have identical losses.

15.3 Complete acquisition results and diagnostic scope

Table 17: All eleven acquisition policies in the executed OCM follow-up.

Acquisition	Hist. F0	Hist. F1	New F0	New F1
Random	3.705	3.092	5.476	4.787
Operating V	3.821	3.300	4.817	4.483
Composition V	3.370	2.708	5.039	4.213
Feature cover	3.419	2.921	4.956	5.299
Pool V	3.657	2.886	5.175	4.015
Kernel V	3.458	2.852	5.010	4.442
Robust multi-view V	3.379	2.702	5.039	4.213
Fixed historical best	3.370	2.708	5.039	4.213
Gate without margin	3.569	2.987	5.039	4.213
Conservative view gate	3.373	2.711	5.039	4.213
Shuffled composition V	3.647	3.039	5.233	4.714

MAE in percentage points of C_2 yield. Frozen adaptive blend, five labels, random query. Hist.: 22 sources, seed-averaged; new: one publication/family, without a population confidence interval.

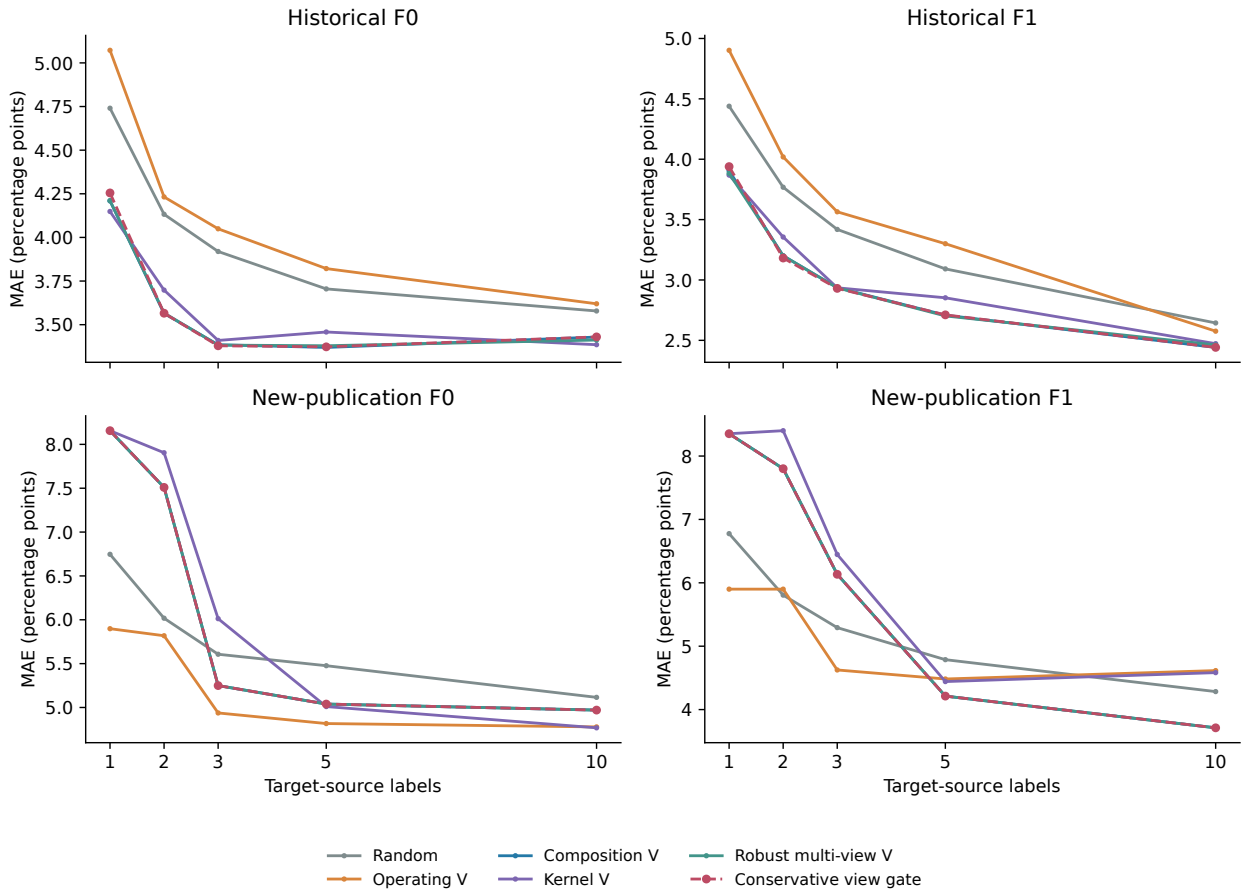


Figure 9: All five label budgets for six central acquisition policies under the fixed adaptive blend. Historical panels average sources equally; new-publication panels describe one fixed source/family. No error bars are created from support draws or from that single family. The conservative gate largely overlaps composition V. Full twelve-predictor cells, source summaries and the remaining policies are included as machine-readable results.

Electronic tables contain all twelve predictor results at five labels:

- `reports/step9_history_all_predictors_k5.csv`
- `reports/step9_new_publication_all_predictors_k5.csv`

The respective `runs/*/aggregated/summary.csv` files contain every policy, predictor, budget and evaluated scenario, including high-temperature query. Historical exploratory paired contrasts use conditional 95% intervals and are explicitly labelled as post-result diagnostics; they are not a new confirmatory family. The five-panel primary results in the preceding sections are retained unchanged. A new random seed convention explains small differences between the historical follow-up's random means and the earlier five-panel means.

16 New measurement curation and frozen application

16.1 Inclusion, units and original experiment identity

The author dataset of Nishimura et al. [22] reports measured temperature curves, validation phases, furnace length and feed ratio. Its 2928 rows include water-summary lines, repeated conditions and controls. We include only explicitly identified SVR-1, SVR-2 and Bayesian-1/2/3 validation phases, whose original study describes new catalyst preparation and testing. Mixed literature/HTS List0, proposed/predicted lists, `Unpublished`, `STD`, and `Refernce/reference` blocks are excluded. The retained population contains 1739 rows, 116 unique metal/support catalyst identities, one publication and one experiment family. The exclusion ledger records 1024 phase exclusions, 120 incomplete temperature/response entries and 45 repeated catalyst conditions. The first CSV occurrence per furnace/feed/slot-invariant composition/support/temperature signature is retained. No query responses are averaged into a support response.

Decimal commas become decimal points and Celsius becomes Kelvin by adding 273.15. Only explicitly reported feed ratios are used. The author protocol reports a 270 mm furnace, 21/7/3 mL min^{-1} $CH_4/O_2/N_2$ and 50 mg catalyst, but these facts do not supply an unambiguous gas residence volume or explicitly resolved pressure; total pressure, reactant partial pressures and contact time remain missing. The missing fractions for all four fields are 1.0. Historical-only medians are used by the frozen transformer. Precursor mmol per support mass is not converted into the historical mol-percent basis. All mol-percent fields remain missing. Single-cation oxide supports map to their support-cation symbols while original formulas are retained; oxidation state or morphology is not inferred.

Reported ethane and ethylene carbon yields sum to C_2 yield within 0.1 percentage points, and methane conversion minus reported carbon products agrees with missing carbon within 0.2 percentage points, consistent with one-decimal rounding and the original N_2 internal-standard procedure. The carbon-basis endpoint is not interpreted as an outlet mole fraction. Labels outside physical bounds would be excluded with a reason, never clipped; no such exclusion was required among eligible complete responses.

Rounded fingerprint screening finds nine old-row matches for three low-yield Na–Mn–W/SiO_2 records: 0.0% at 600/650 °C and 0.5% at 700 °C, versus old 2011 records of 0.0/0.6%. The new paper explicitly includes that catalyst among its newly prepared SVR-1 validation measurements. These nonspecific low-yield coincidences do not establish copied experimental runs; the full audit and review notes are retained. Original measurement origin, phase tracing and exclusion of mixed/reused inputs support new publication-level application. They do not certify all historical laboratory identities or full catalyst-batch independence. All phases are therefore one inference group.

16.2 Chronology, freezing and implementation verification

The twenty-five historical source/fold comparisons were completed and audited before full-history fitting. The final pipeline fits 54 group-OOF backbone candidates and two full backbones, selects

controls using historical responses only, and freezes 79 code/model/parameter/input hashes before raw candidate download and new-cohort scoring. Public paper conclusions and the CADS landing page had informed screening before this stage; the source is not a prospectively collected campaign or public preregistered test. The included row outcomes were not used to select the gate or to retune the frozen controls. Curation is locked before the first numerical application.

Historical execution exports 93150 cells, 662400 query-row records and 7948800 scalar predictions. All scalar losses are independently reaggregated, all 9315 support orders reconstructed, and 105600 query-row predictions regenerated through the prediction helpers. Actual outcome poisoning tests 2040 source/split frames, including 110 eligible frames, comparing 270 eligible partition index arrays. Separate numerical checks validate dual/primal variance equality, ridge query and exact leave-one-out values against direct scikit-learn fits, and equivalent-view robust/linear order equality. The published-source application exports 1380 cells, 620310 query-row records and 7443720 scalar values; all scalar losses and 138 support orders pass auditing and 98890 query-row predictions are reconstructed. The reconstruction reuses prediction helpers; only the loss reaggregation and direct math references are independent implementations. None of these checks is a new biological/chemical replication or a pipeline-refit bootstrap.

The ACS 2024 SI candidate was downloaded and its publisher-provided MD5 verified. Its experimental Figure S8 provides plotted replicates for three catalysts, without a complete compatible numeric row table in that asset; it is not silently digitized and counted as a second new dataset. The additional methanol candidate remains unevaluated because its publisher landing access returns HTTP403 and no compatible SI rows were acquired. These sources contribute no extra source count or numerical performance claim.

16.3 Portable reproduction and opened-cohort boundary

The new root `protocol/runtime_authorization.json` records the user's explicit replacement of the earlier private-server-only constraint. Entry points permit the authorized current runtime and keep outputs inside the project. `server/run_history.sh` performs the math check, real one-part smoke, full historical grid, audit, aggregation and final freezing. Archived Step7 bases and OOF responses are inputs to the historical extension, not claimed new outer refits. The original prepared Step8 guard/instructions are retained as predecessor provenance.

`analysis/curate_nishimura.py` writes all inclusion, provenance, missing-unit and overlap files; `code/evaluate_external_ocm.py` verifies the freeze hashes and applies the same fixed pipeline once. The manuscript sources and `analysis/build_step9_results.py` regenerate the measured figures/tables without fitting. The completed new cohort is now opened evidence. Any further gate redesign must treat it as development and obtain a different untouched publication/laboratory or a prospectively locked experiment for confirmation.

References

[1] Toldy, Á. I. et al. TheMeCat: Dataset of Thermocatalytic Conversion of CO_2 to Methanol. *Scientific Data* **13**, 165 (2026). https://doi.org/10.1038/s41597-025-06482-8.

[2] Yan, M. et al. A machine learning perspective on three decades of methanol synthesis: research framework and experimental operation insights. *Communications Engineering* **5**, 162 (2026). https://doi.org/10.1038/s44172-026-00706-4.

[3] Suvarna, M., Araújo, T. P. & Pérez-Ramírez, J. A generalized machine learning framework to predict the space-time yield of methanol from thermocatalytic CO_2 hydrogenation. *Applied Catalysis B: Environmental* **315**, 121530 (2022). https://doi.org/10.1016/j.apcatb.2022.121530.

[4] Schmack, R. et al. A meta-analysis of catalytic literature data reveals property–performance correlations for the OCM reaction. *Nature Communications* **10**, 441 (2019). https://doi.org/10.1038/s41467-019-08325-8.

[5] Mine, S. et al. Analysis of Updated Literature Data up to 2019 on the Oxidative Coupling of Methane Using an Extrapolative Machine-Learning Method to Identify Novel Catalysts. *ChemCatChem* (2021). https://doi.org/10.1002/cctc.202100495. Data: https://github.com/mts-uw/OCM, commit dbf7be5aadb30323f825e8f7d7f6bd3012c99dd2.

[6] Roberts, D. R. et al. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* **40**, 913–929 (2017). https://doi.org/10.1111/ecog.02881.

[7] Wu, W., Yang, M., Cheng, P. & Wang, Y. Reliable Machine Learning Screening of Adsorption Energies Is Better Assessed with Formula-Grouped Cross-Validation. *Catalysts* **16**, 742 (2026). https://doi.org/10.3390/catal16080742.

[8] Ben-David, S. et al. A theory of learning from different domains. *Machine Learning* **79**, 151–175 (2010). https://doi.org/10.1007/s10994-009-5152-4.

[9] Feudale, R. N. et al. Transfer of multivariate calibration models: a review. *Chemometrics and Intelligent Laboratory Systems* **64**, 181–192 (2002). https://doi.org/10.1016/S0169-7439(02)00085-0.

[10] Efron, B. & Morris, C. Stein's estimation rule and its competitors—an empirical Bayes approach. *Journal of the American Statistical Association* **68**, 117–130 (1973).

[11] Breiman, L. Stacked Regressions. *Machine Learning* **24**, 49–64 (1996). https://doi.org/10.1023/A:1018046112532.

[12] Xu, W. H. et al. A natural language processing to causality framework for robust knowledge extraction of CO_2 hydrogenation with batch effect control. *npj Computational Materials* **12**, 191 (2026). https://doi.org/10.1038/s41524-026-02065-2.

[13] K. Yu, J. Bi and V. Tresp. Active learning via transductive experimental design. ICML, 1081–1088 (2006). doi:10.1145/1143844.1143980.

[14] X. Fontaine, P. Perrault, M. Valko and V. Perchet. Online A-optimal design and active linear regression. ICML, PMLR 139, 3374–3383 (2021). https://proceedings.mlr.press/v139/fontaine21a.html.

[15] S. J. Sloman, A. Bharti, J. Martinelli and S. Kaski. Bayesian active learning in the presence of nuisance parameters. UAI, PMLR 244, 3245–3263 (2024). https://proceedings.mlr.press/v244/sloman24a.html.

[16] G. Hooker. Generalized functional ANOVA diagnostics for high-dimensional functions of dependent variables. J. Comput. Graph. Stat. 16, 709–732 (2007). doi:10.1198/106186007X237892.

[17] C. E. Rasmussen and C. K. I. Williams. Gaussian Processes for Machine Learning. MIT Press (2006), Chapter 2, Section 2.7. https://gaussianprocess.org/gpml/chapters/RW2.pdf.

[18] J. Harrison, A. Sharma and M. Pavone. Meta-Learning Priors for Efficient Online Bayesian Regression. arXiv:1807.08912 (2018). https://arxiv.org/abs/1807.08912.

[19] Q. P. Nguyen, B. K. H. Low and P. Jaillet. Learning to learn with Gaussian processes. UAI, PMLR 161, 1466–1475 (2021). https://proceedings.mlr.press/v161/nguyen21c.html.

[20] P. Tighineanu et al. Scalable Meta-Learning with Gaussian Processes. AISTATS, PMLR 238, 1981–1989 (2024). https://proceedings.mlr.press/v238/tighineanu24a.html.

[21] T. Iwata and A. Kumagai. Meta-learning to Calibrate Gaussian Processes with Deep Kernels for Regression Uncertainty Estimation. arXiv:2312.07952 (2023). https://arxiv.org/abs/2312.07952.

[22] S. Nishimura, X. Li, J. Ohyama and K. Takahashi. Leveraging machine learning engineering to uncover insights into heterogeneous catalyst design for oxidative coupling of methane. Catal. Sci. Technol. 13, 4646–4655 (2023). doi:10.1039/d3cy00596h. Author measurements: https://cads.eng.hokudai.ac.jp/datamanagement/datasources/d84c1e22-ceb9-488a-8d45-4c7cf1c603b5/.

[23] A. Gorodetsky and Y. Marzouk. Mercer Kernels and Integrated Variance Experimental Design: Connections Between Gaussian Process Regression and Polynomial Approximation. SIAM/ASA J. Uncertain. Quantif. 4, 796–828 (2016). doi:10.1137/15M1017119.